

Search

COMP9414 24t2

version 1.1

Dr Armin Chitizadeh



UNSW
SYDNEY

By the end of this lecture you will learn

- ❖ The **theory** behind **search in AI**
- ❖ To develop a **smart agent** with **search capability** to solve complicated **problem**



UNSW
SYDNEY

Content

- What is Search and why is it needed?
- Human Search vs Computer Search
- Converting real-world problem into machine understandable problem
- Choosing the best search algorithm



UNSW
SYDNEY

What is Search and Why is it needed

“

In which we see how an agent can find a sequence of actions that achieves its goals when no single action will do.

- Peter Norvig and Stuart J. Russell
Artificial Intelligence: A Modern Approach

”



UNSW
SYDNEY

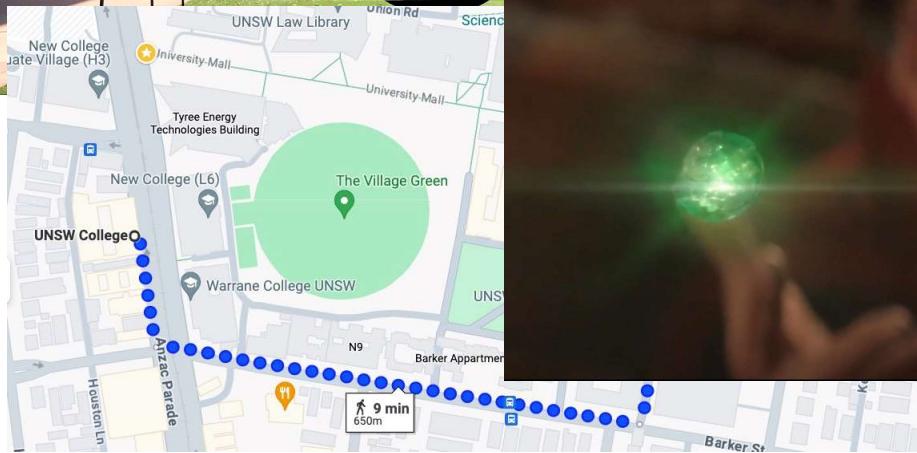
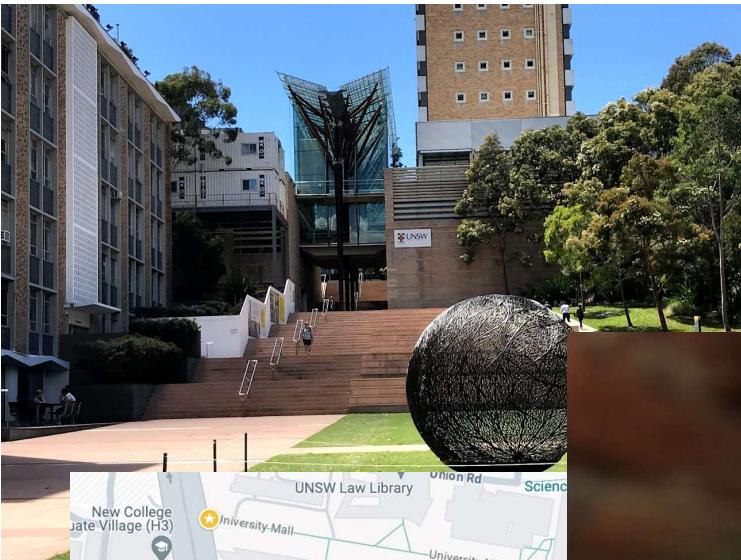


Photo by Charlie Solorzano
from pexels.com



When is Search Needed?

Motion Planning

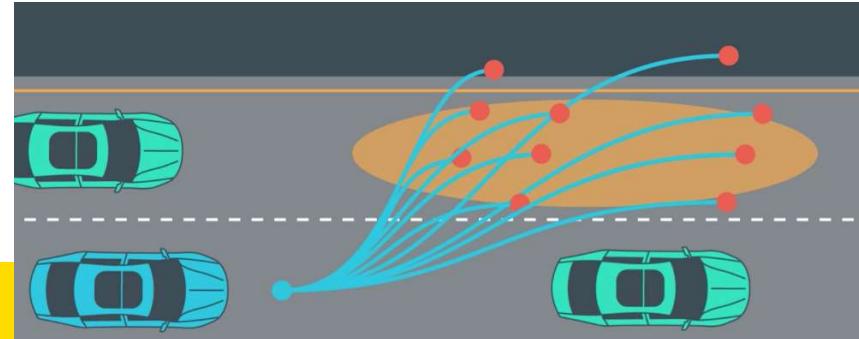
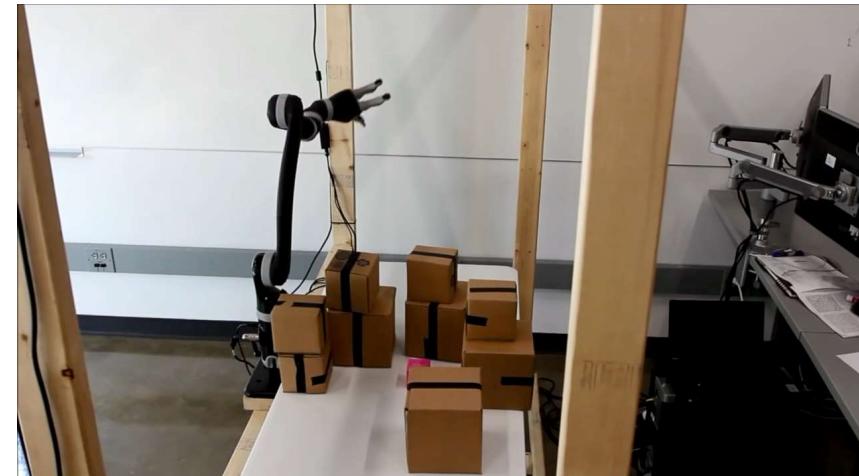
Navigation

Speech and Natural Language

Task Planning

Machine Learning

Game Playing



Human Search vs Computer Search



UNSW
SYDNEY

Human Search vs Computer Search



Photo generated by Microsoft Copilot



UNSW
SYDNEY

Human Search vs Computer Search



Photo generated by Microsoft Copilot



UNSW
SYDNEY

Converting real-world problem into machine understandable problem



UNSW
SYDNEY

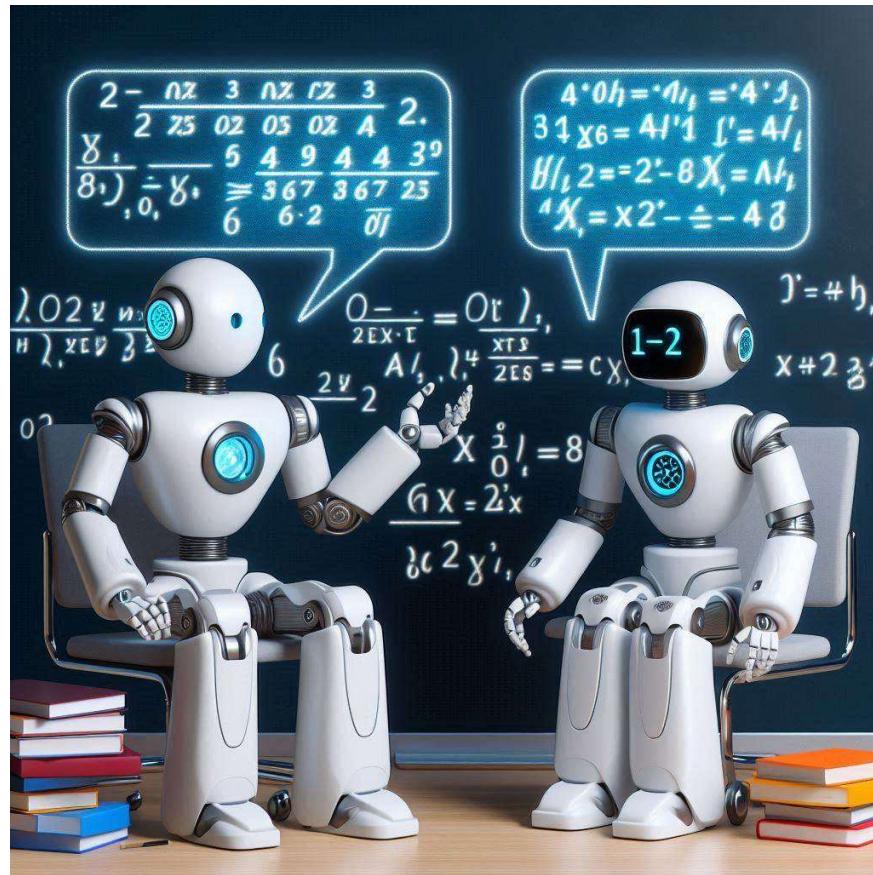


Photo generated by Microsoft Copilot



UNSW
SYDNEY

Let's formalize it!

A problem can be described by

- state

2	3	5
8	6	1
7	4	

$[[2,3,5], [8,6,1], [7,4,0]]$



$[[rB, 0, bB, qB, KB, bB, 0, 0], [pB, 0, pB, pB, 0, 0, 0, 0], [0, 0, kB, 0, 0, pW, 0, rB], \dots]$

- action

{"up", "left"}

{ a5 , a3 , ... }

- Transition Function do: $S \times A \rightarrow S$



UNSW
SYDNEY

Let's formalize it

Transition Function

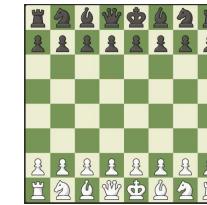
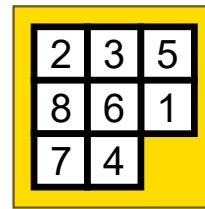


UNSW
SYDNEY

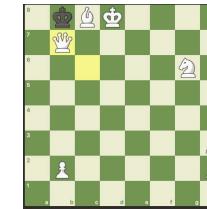
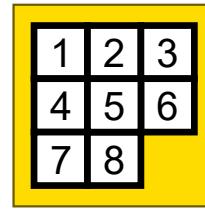
Let's formalize it

Special states:

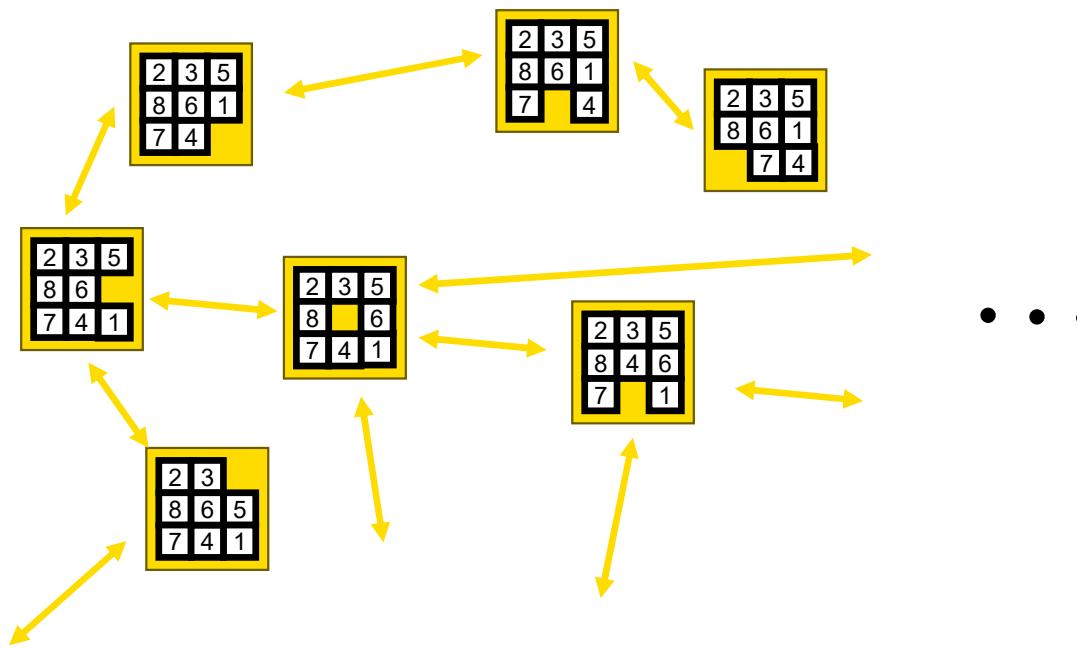
- Initial state: s_0



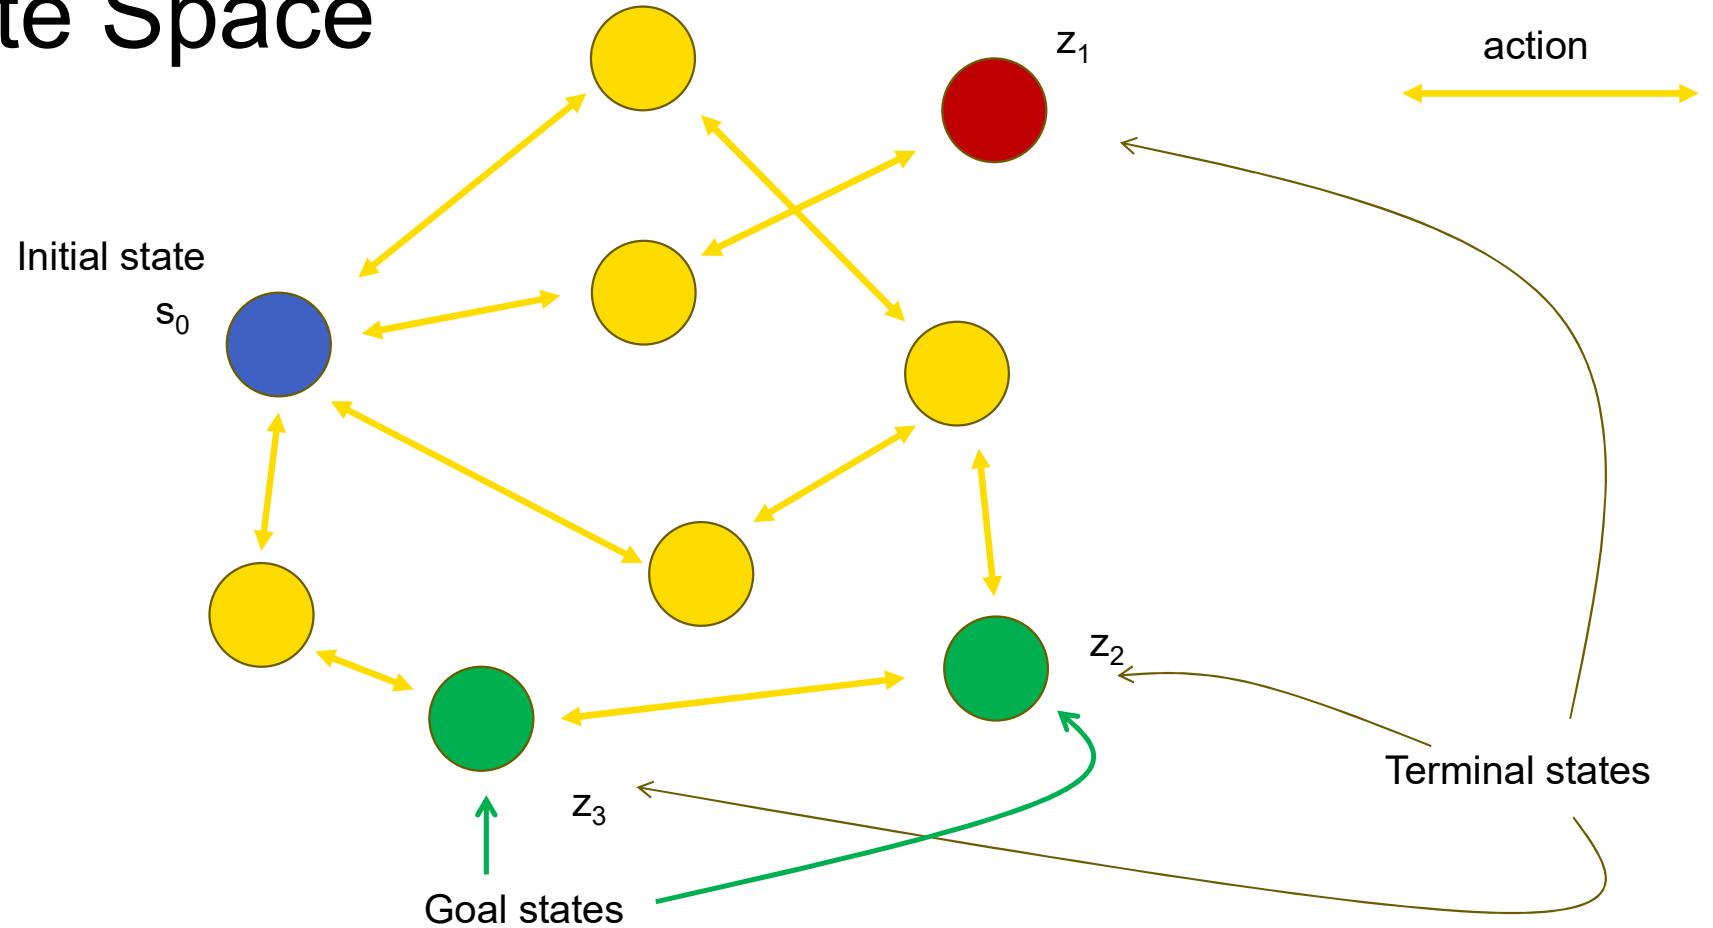
- Terminal states: $Z \subset S$



State Space



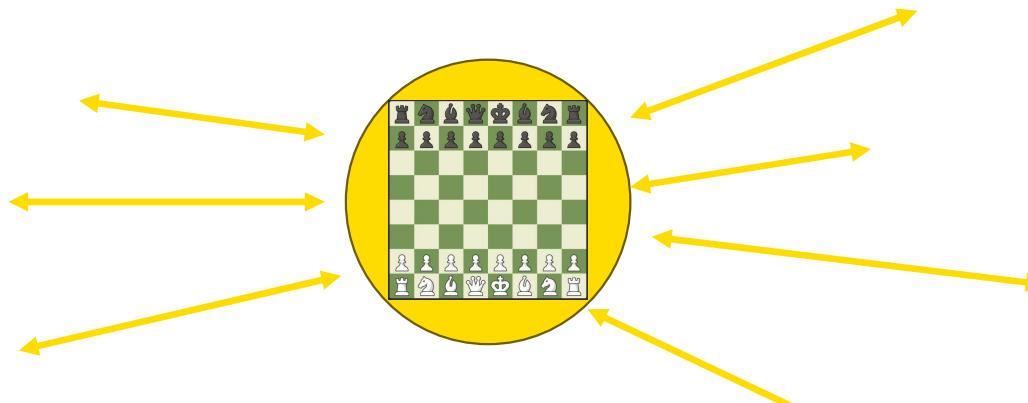
State Space



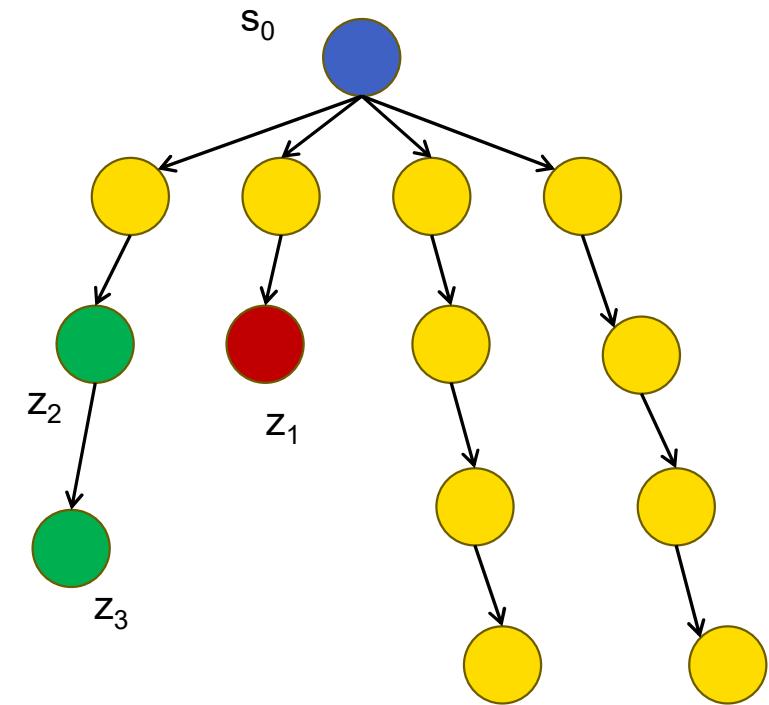
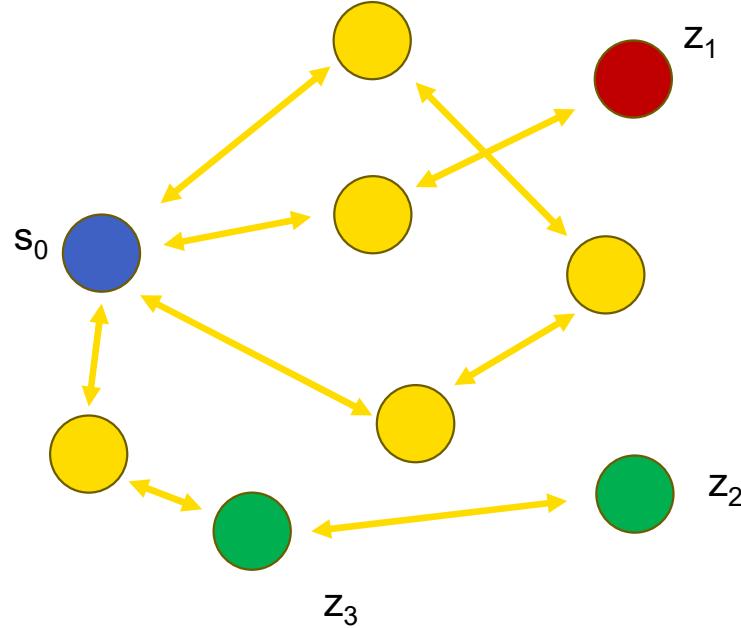
State Space in Graph

Node vs State

- Node in Graph is a data structure which contains a state and other related information such as legal actions from the state



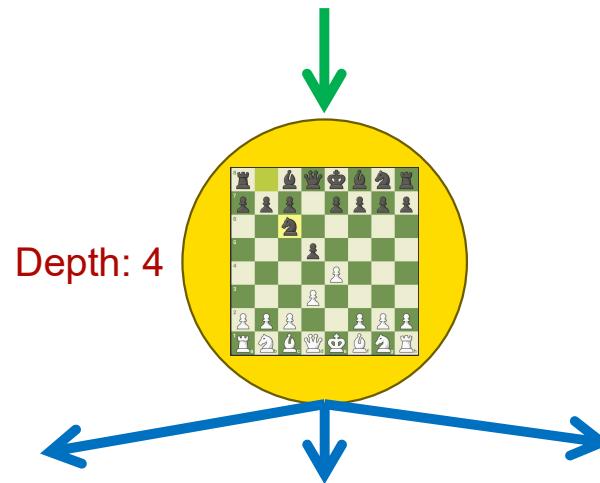
Graph State Space vs Tree State Space



State Space in Tree

Node vs State

- Node in Tree is a data structure which contains a state and other related information such as **children nodes**, **parent node** and **depth**



Now we have converted problems in their tree structure, can we solve all problems?

Approximately 10^{24} possible states



Photo by Charlie Solorzano
from pexels.com

- For many problems you can't generate the whole tree
- So you can not search all possible path
- When you have a choice, which path will you choose?



UNSW
SYDNEY

**Any Question on Comp9414
Armin's Lecture**



<https://forms.office.com/r/bBb6bt2sv7>



UNSW
SYDNEY

Choosing the best Search algorithm



UNSW
SYDNEY

Search Algorithms Covered at this Lecture	
Uninformed Search	Breath First Search (BFS)
	Depth First Search (DFS)
	Depth Limited Search
	Bidirectional Search
	Iterative Deepening Search (IDDFS)
	Uniform Cost Search (UCS)
Informed Searches	<u>Heuristic</u>
	Greedy Best-First Search
	A* Search



Overview Uninformed Search

Informed Searches	Background & Terminology
	Breath First Search (BFS)
	Depth First Search (DFS)
	Uniform Cost Search (UCS)
	Depth Limited Search
	Iterative Deepening Search (IDDFS)
	Bidirectional Search
	Heuristic
Uninformed Search	Greedy Best-First Search
	A* Search



Search Methods

Uninformed search

- use no problem-specific information
- Uninformed (or “blind”) search strategies use only the information available in the problem definition (can only distinguish a goal from a non-goal state)

Informed search



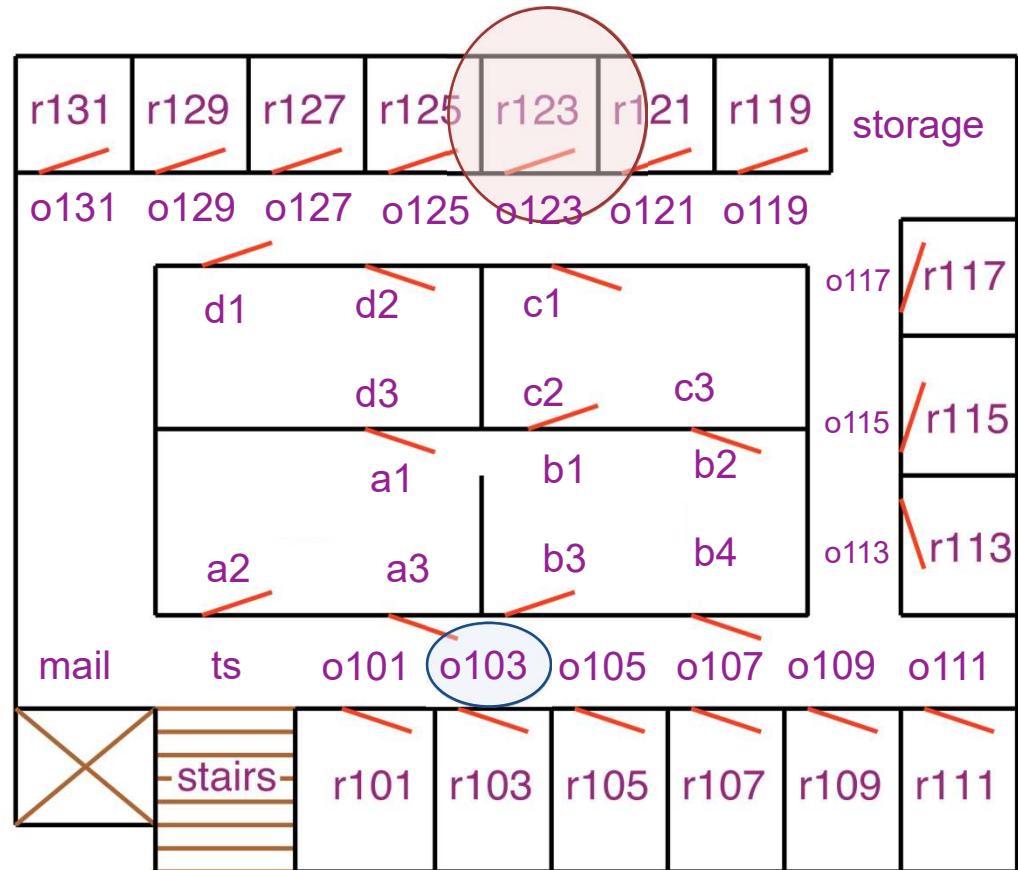
UNSW
SYDNEY

Delivery Robot

The robot wants to get from outside room 103 to the inside of room 123.

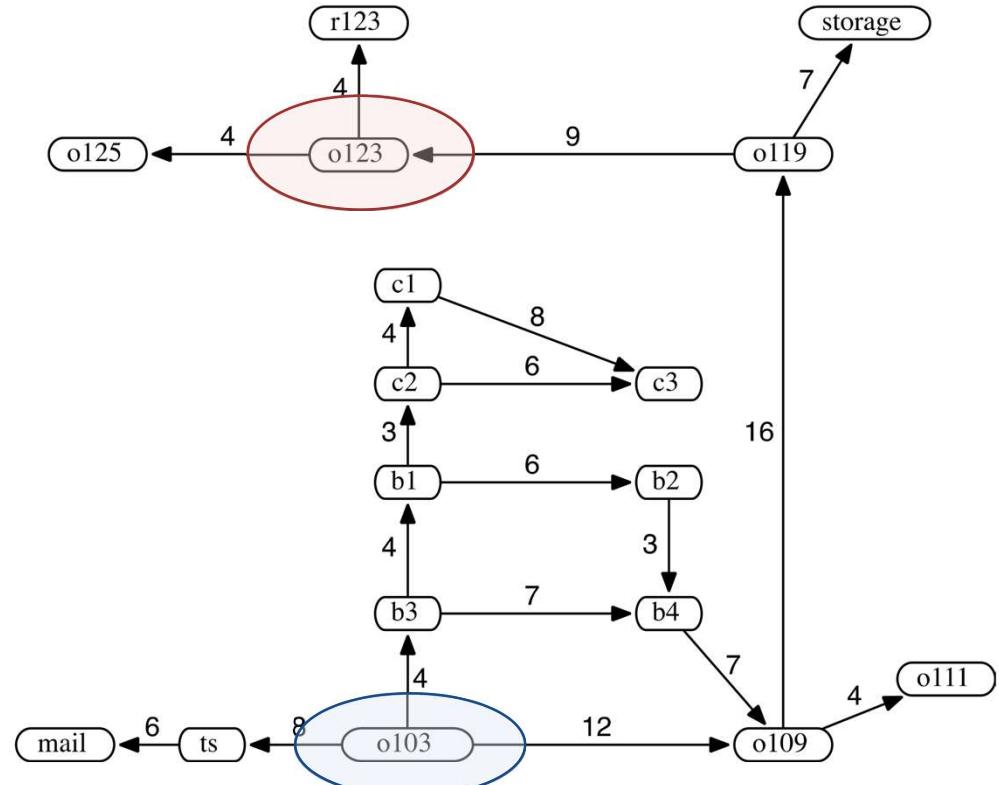
- The only way a robot can get through a doorway is to push the door open in the direction shown.

The task is to find a path from o103 to r123



State-Space Graph for Delivery Robot

Modelled as a state-space search problem
States are locations.



State Space Search Problems

State space — set of all states reachable from initial state(s) by any action sequence

Initial state(s) — element(s) of the state space

Transitions

- **Operators** — set of possible actions at agent's disposal; describe state reached after performing action in current state, or
- **Successor function** — $s(x)$ = set of states reachable from state x by performing a single action

Goal state(s) — element(s) of the state space

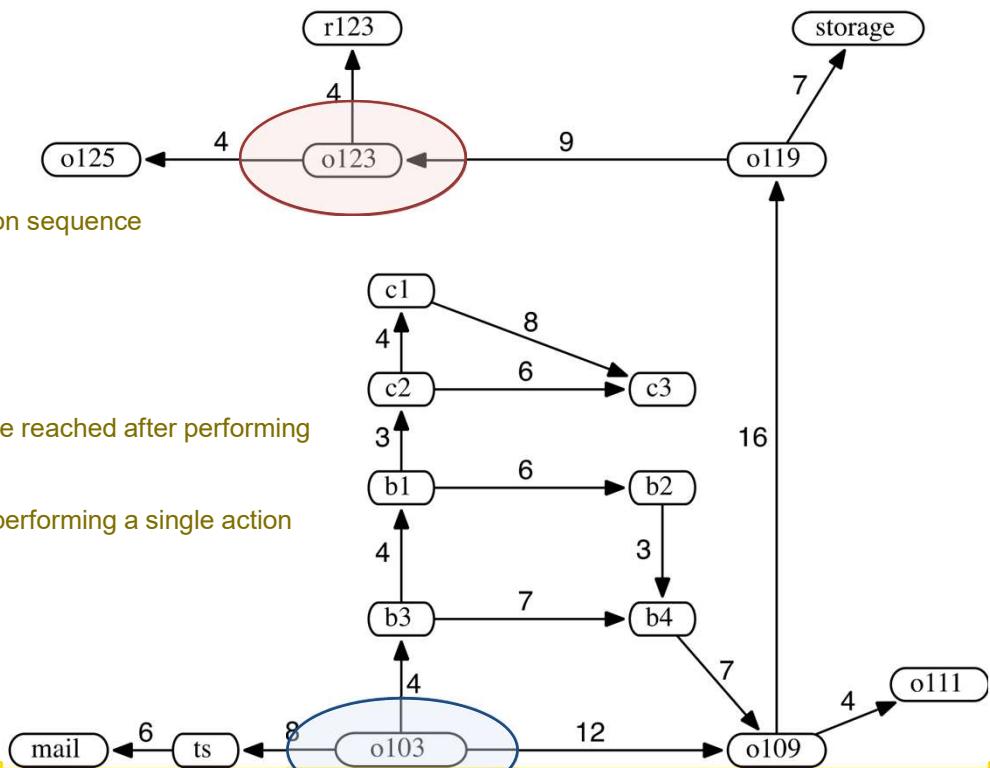
Path cost — cost of a sequence of transitions used to evaluate solutions (applies to optimisation problems)



UNSW
SYDNEY

State-Space Graph for Delivery Robot

- Modelled as a state-space search problem
- States are locations.
- **State space** — set of all states reachable from initial state(s) by any action sequence
- **Initial state(s)** — element(s) of the state space
- Transitions
 - **Operators** — set of possible actions at agent's disposal; describe state reached after performing action in current state, or
 - **Successor function** — $s(x)=$ set of states reachable from state x by performing a single action
- **Goal state(s)** — element(s) of the state space
- **Path cost** — cost of a sequence of transitions used to evaluate solutions
(applies to optimisation problems)

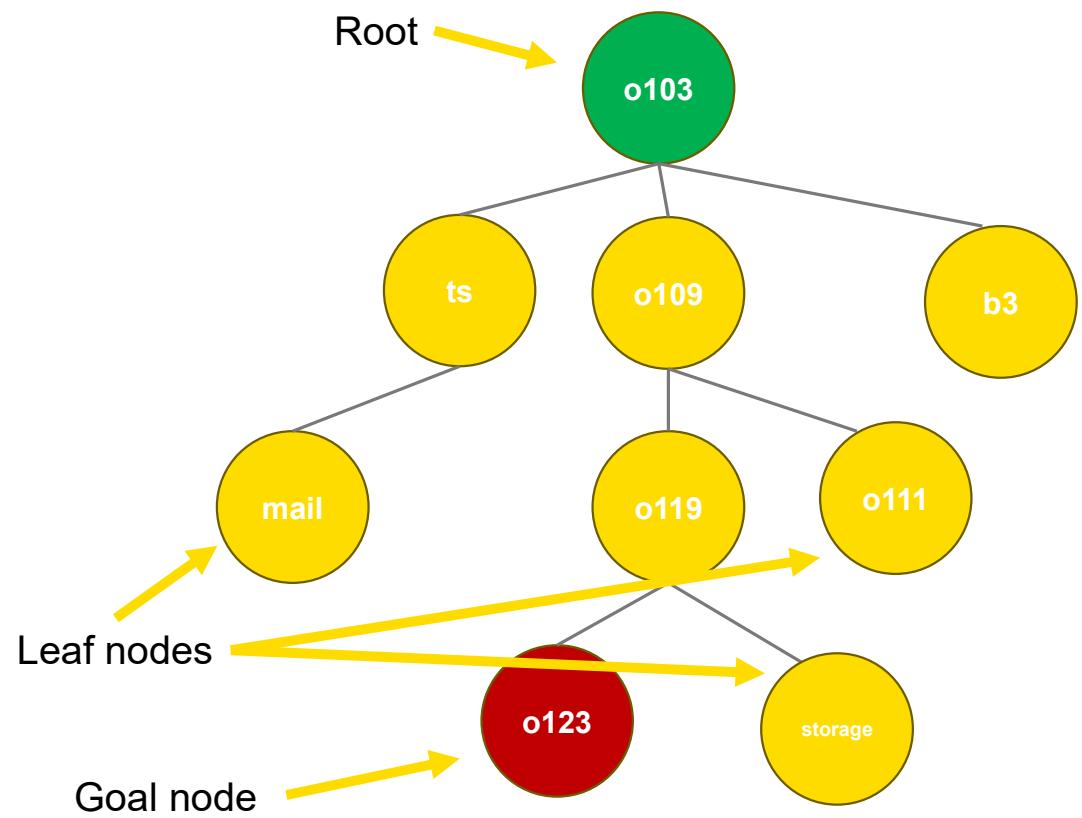


Search Tree

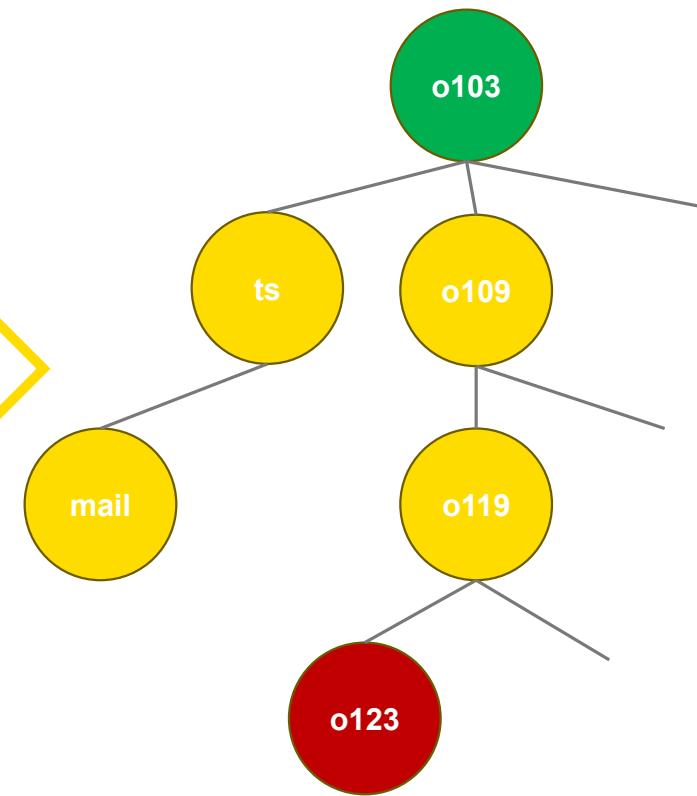
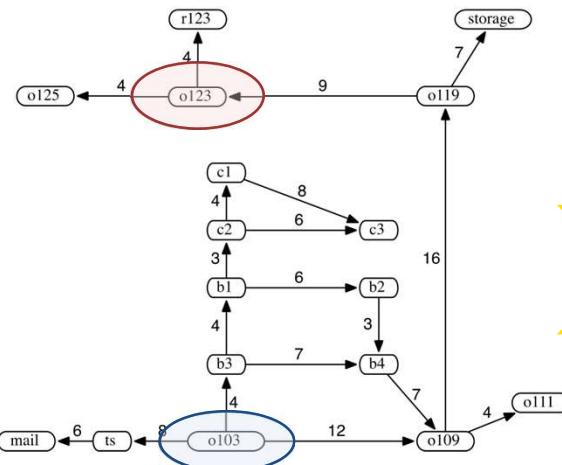
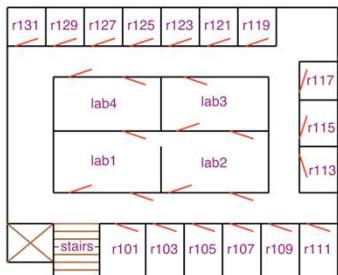
Search tree: superimposed over the state space.

Root: search node corresponding to the initial state.

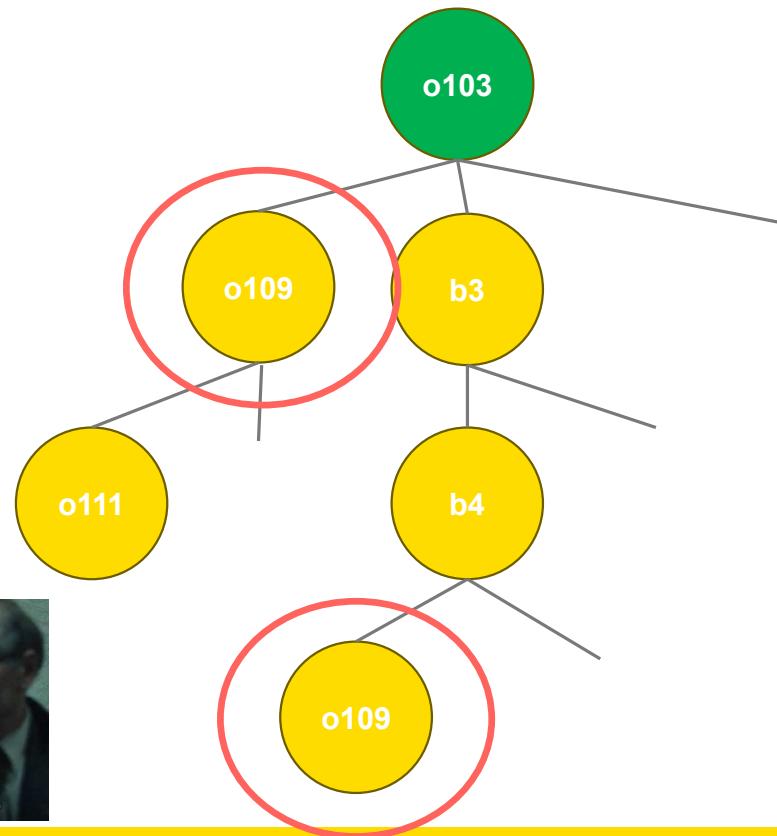
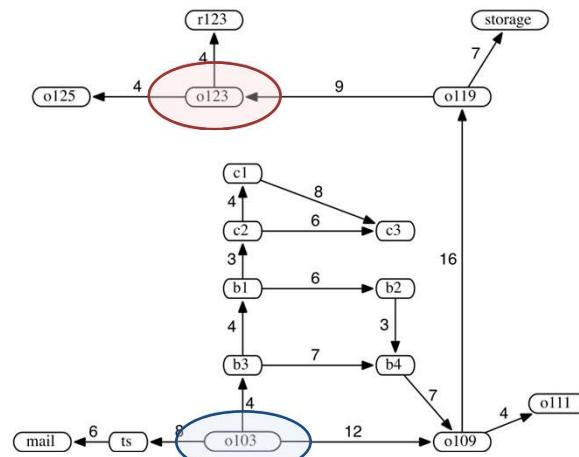
Leaf nodes: correspond to nodes that have no successors in the tree because do not have any successor or are not expanded



Problem to State-Space to Search Tree



Problem to State-Space to Search Tree



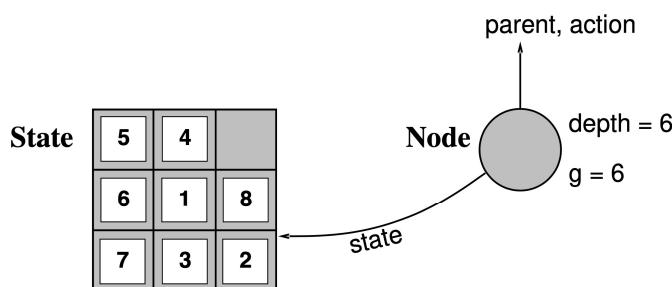
What about Chess?



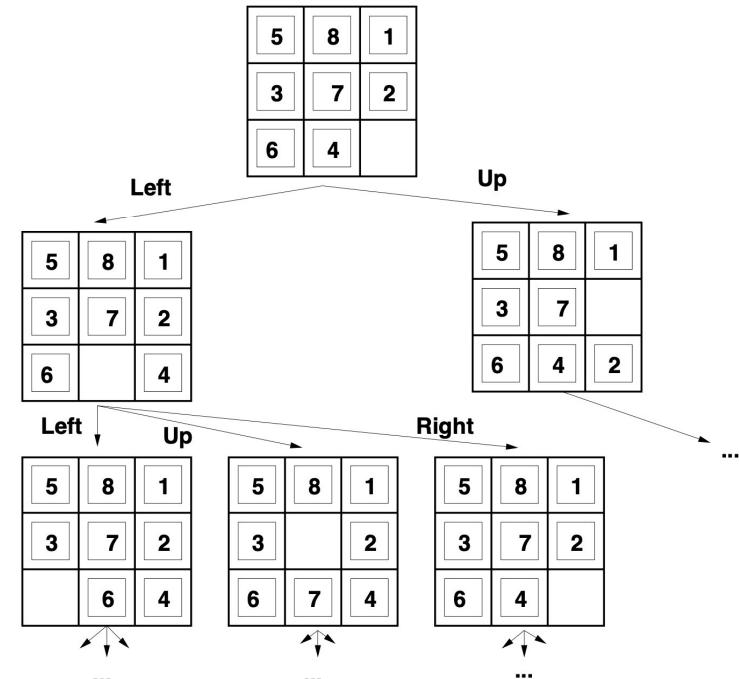
UNSW
SYDNEY

State-Space Graph

- **State** is a representation of problem at a given time.
- In **State-Space Graph or Tree Search** each node contains a states and arcs are actions



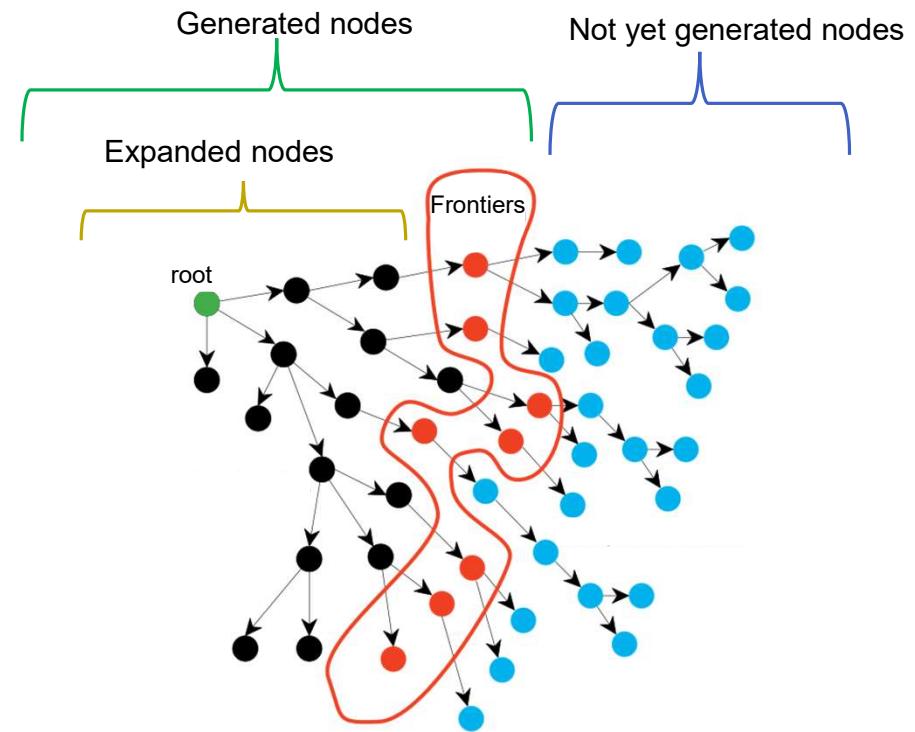
States do not have parents, children, depth or path cost!



Problem Solving by Graph Searching

- **Expanded nodes:** green and black
- **Frontiers:** red
- Generated: green, black and red
- Not yet generated nodes: blue

Search strategy differ in the way they expand the frontier

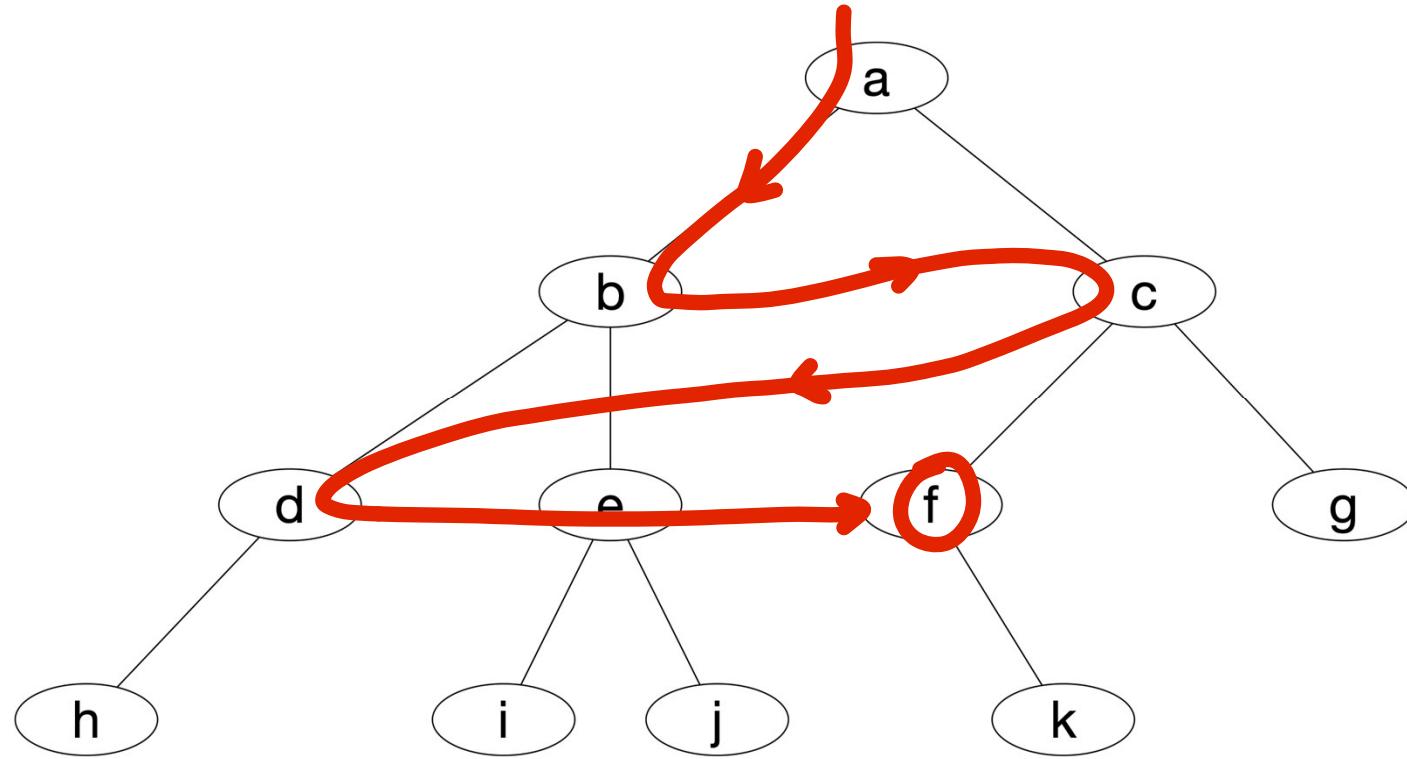


Breath First Search (BFS)



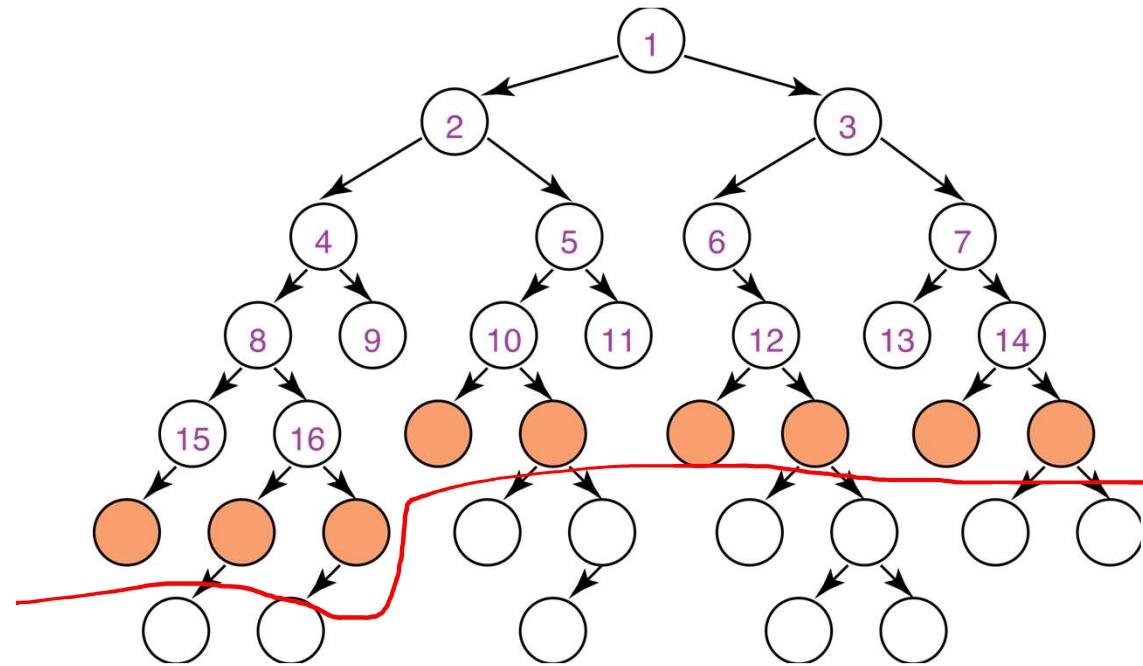
UNSW
SYDNEY

Breadth-First Search

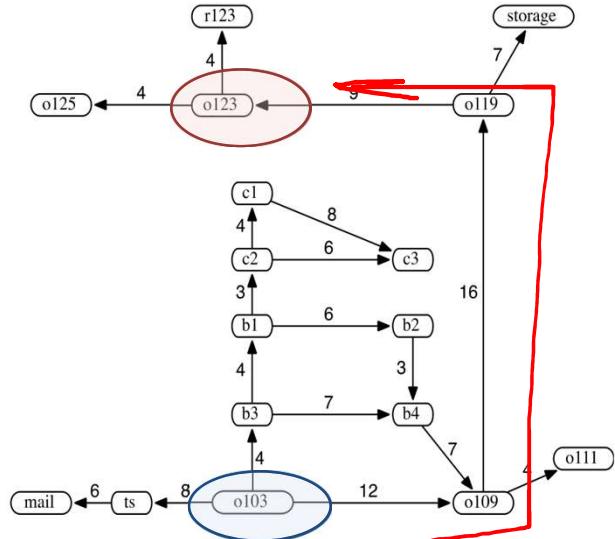
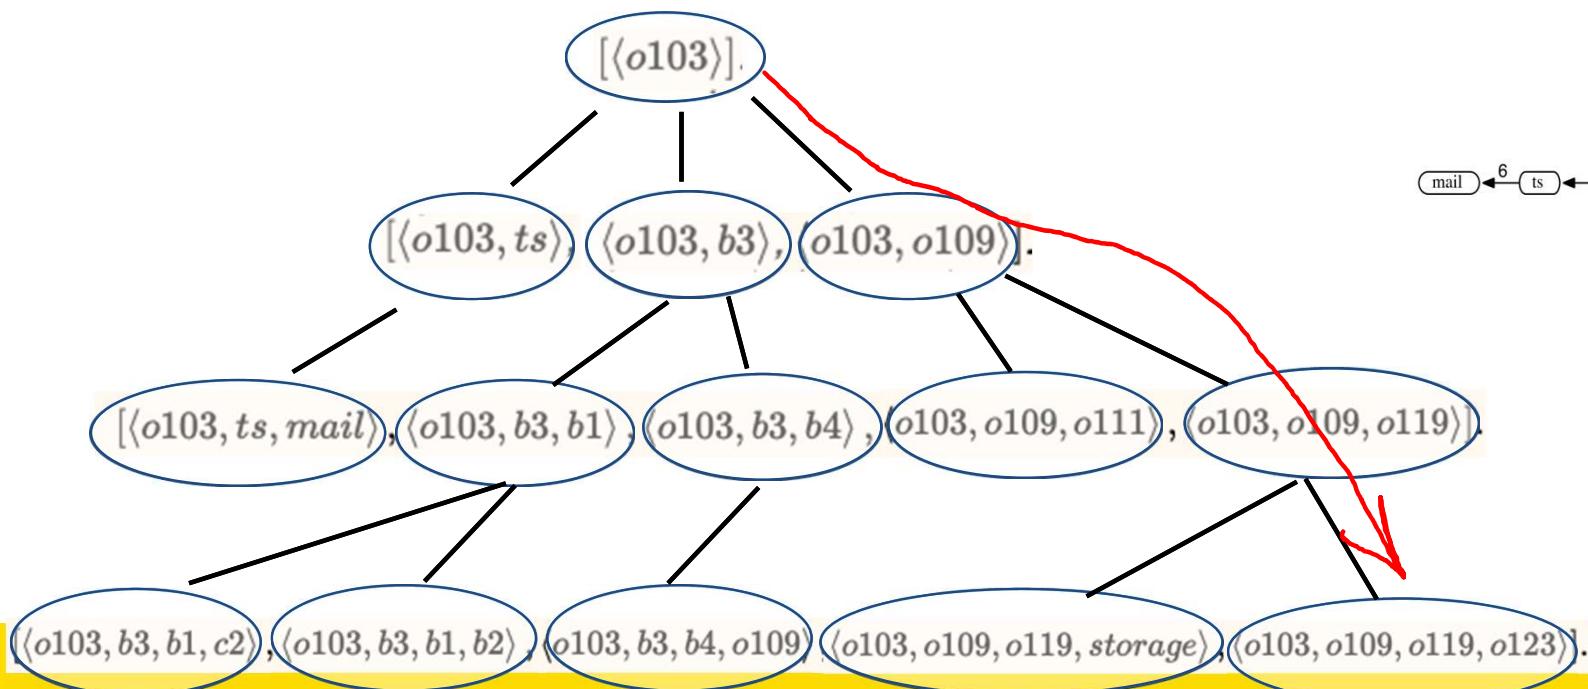


UNSW
SYDNEY

Breadth-first Search Frontier



Breadth-First Search



Breadth-first Search

Breadth-first search treats the frontier as a queue

It selects the first element in the queue to explore next

If the list of paths on the frontier is $[p_1, p_2, \dots, p_r]$:

- p_1 is selected. Its children are added to the end of the queue, after p_r .
- p_2 is selected next.



UNSW
SYDNEY

Breadth-First Search

All nodes are expanded at a same depth in the tree before any nodes at the next level are expanded

Can be implemented by using a queue to store frontier nodes

- put newly generated successors at end of queue

Include check that state has not already been explored

- Needs a new data structure for set of explored states

Finds the shallowest goal first



UNSW
SYDNEY

Complexity of Breadth-first Search

- Does breadth-first search guarantee finding the shortest path?
- What happens on infinite graphs or on graphs with cycles if there is a solution?
- What is the time complexity as a function of the length of the path selected?
- What is the space complexity as a function of the length of the path selected?
- How does the goal affect the search?



UNSW
SYDNEY

Properties of breadth-first search

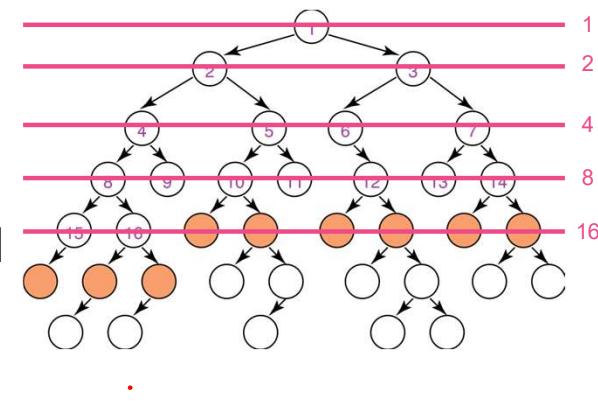
Complete? Yes (if breadth, b , is finite, the shallowest goal is at a fixed depth, d , and will be found before any deeper nodes are generated)

Time? $1 + b^1 + b^2 + b^3 + \dots + b^d = \frac{b^{d+1} - 1}{b - 1} = O(b^d)$

Space? $O(b^d)$ keeps every node in memory; generate all nodes up to level d

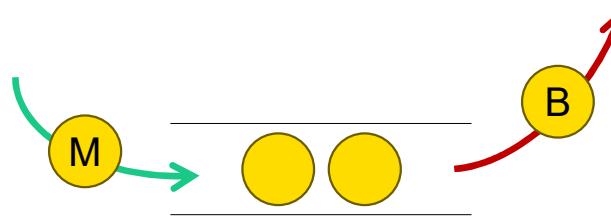
Optimal? Yes, but only if all actions have the same cost

Space is the big problem for BFS. It grows **exponentially** with depth



Expanded vs Generated

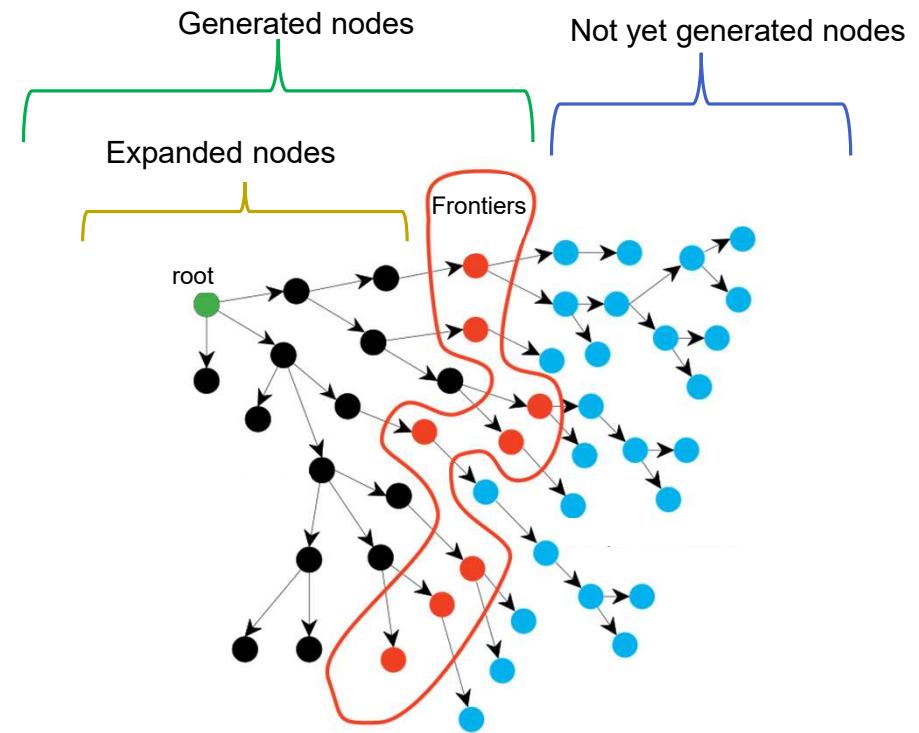
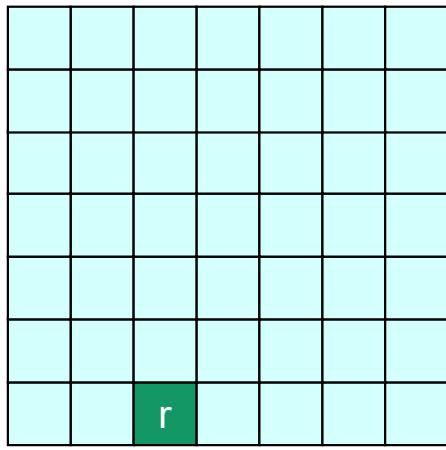
- When a node is inserted in the list, it is **Generated**



Node M is generated
and
Node B is in the process of
being expanded

- When a node is removed from the list, it is in the process of being **Expanded**

BFS in searching a grid



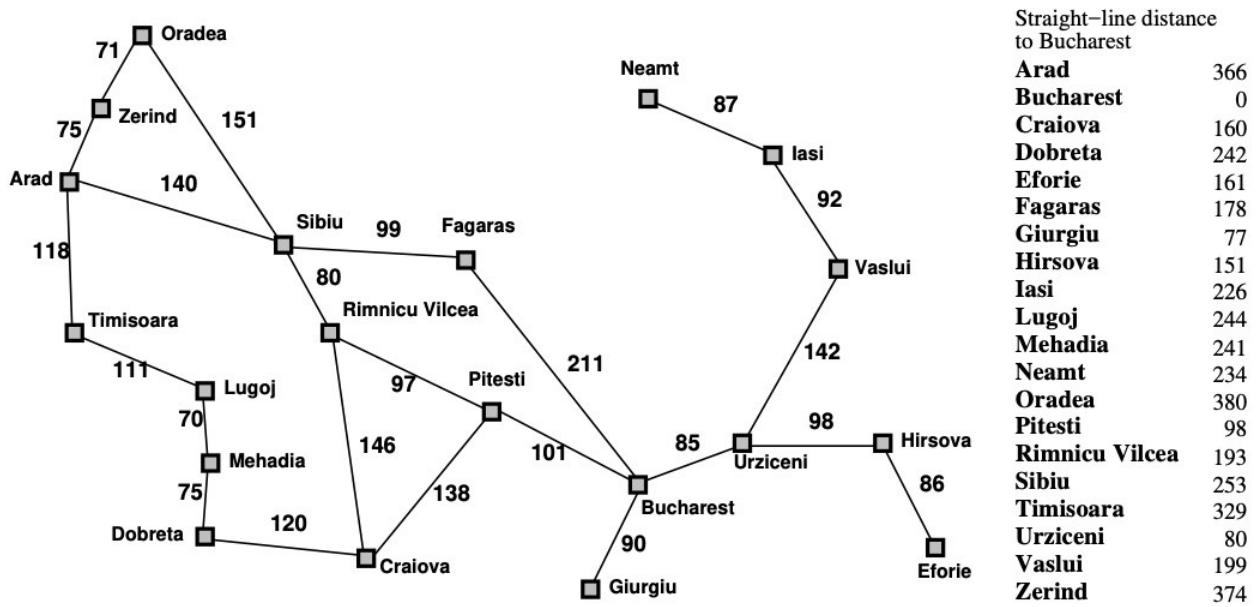
UNSW
SYDNEY

Exercise

This exercise uses the route-finding example with the Romanian map from Russell & Norvig (Artificial Intelligence: A Modern Approach) as an example.

For the route from Arad to Bucharest, what order are nodes in the state space expanded for each of the following algorithms when searching for the shortest path between Arad and Bucharest? Where there is a choice of nodes take the first one by alphabetical ordering. For breath-first search, stop the search when the goal state is generated and use a check to ensure that nodes with the same state as a previously expanded node are not added to the frontier.

* Breadth-First Search

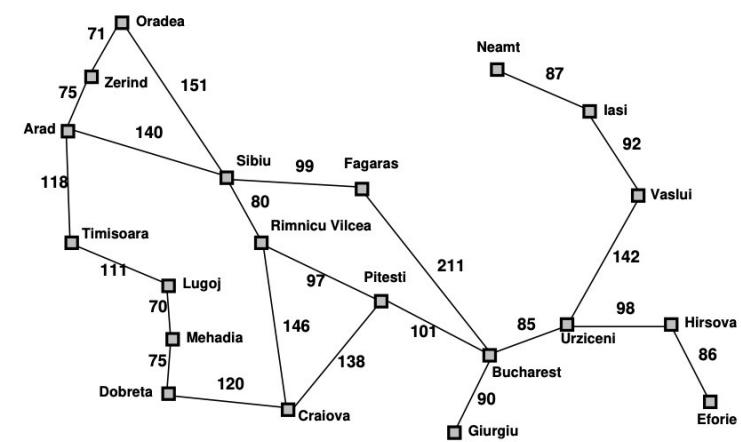


Exercise BFS

B O L R O F Z T S A

expanded

A S T Z F

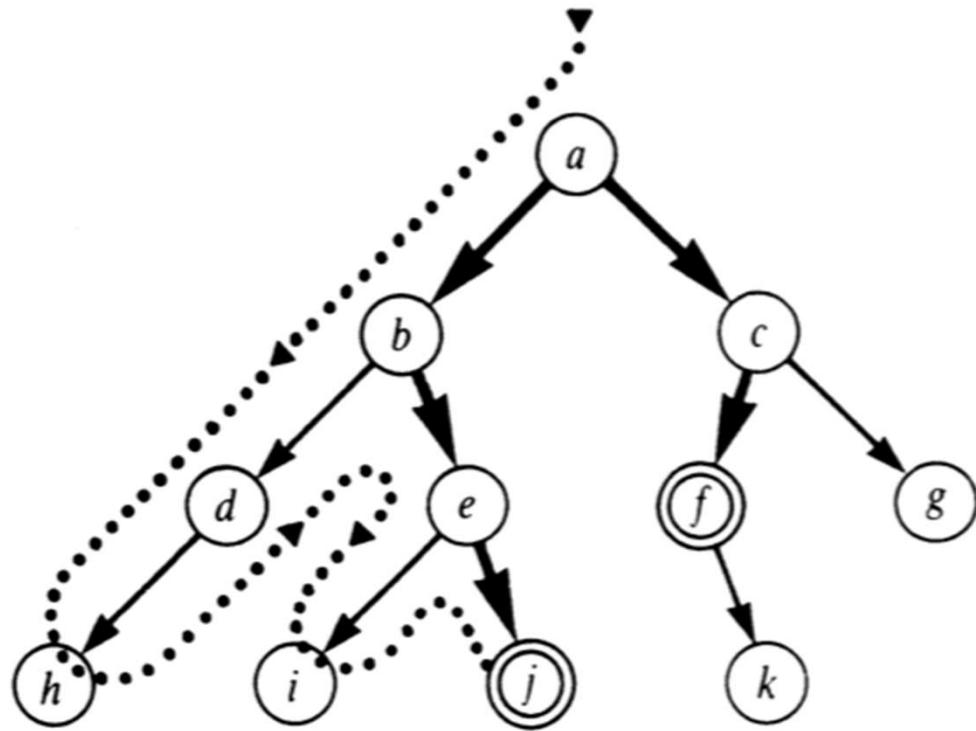


Depth First Search (DFS)



UNSW
SYDNEY

Depth-first Search - DFS



Depth First Search

Expand one node at the deepest level reached so far

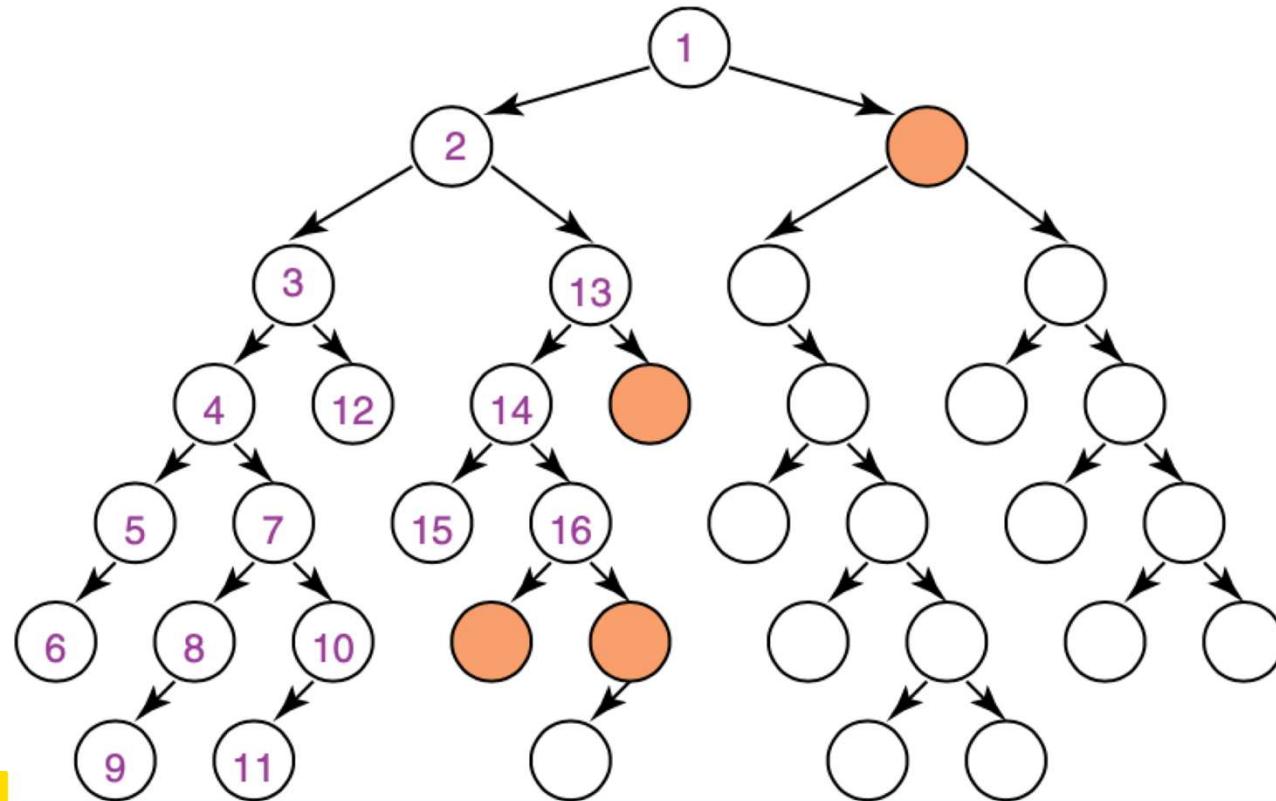
Implementation:

- Implement the frontier as a stack, i.e. insert newly generated states at the front of the open list (frontier)
- Can be implemented by recursive function calls ✓

In depth-first search, like breadth-first, the order in which the paths are expanded does not depend on the goal.

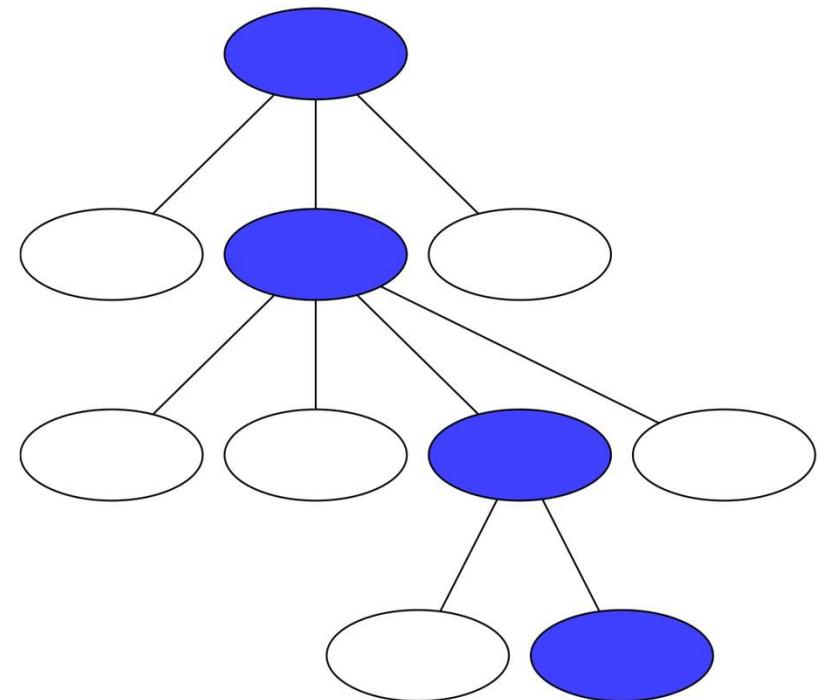


Depth-first Search Example



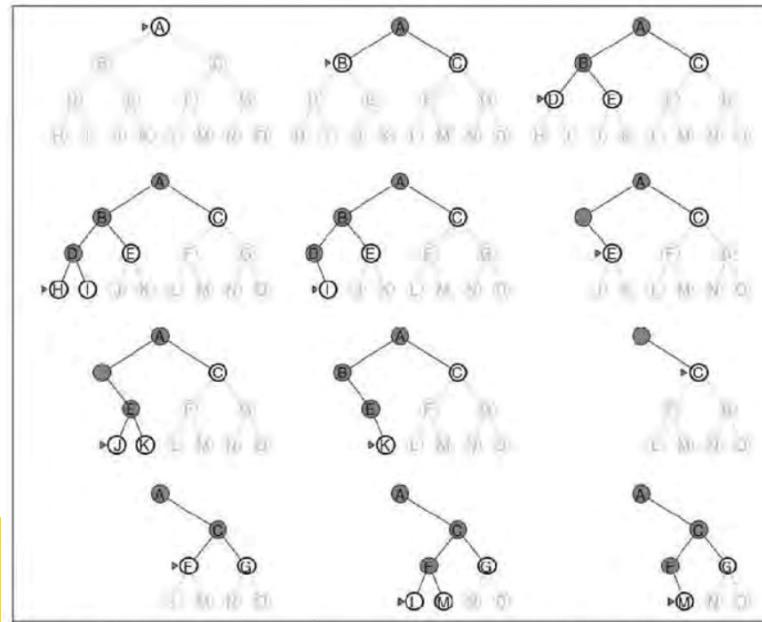
Depth First Search

- At any point depth-first search stores single path from root to leaf, together with any remaining unexpanded siblings of nodes along path
- Stop when node with goal state is expanded
- Include check that state has not already been explored **along a path – cycle checking**



Why DFS is memory efficient?

- Once a node has been expanded, it can be removed from memory as soon as all its descendants have been fully explored. Thus, giving us a time complexity of $O(bm)$



Tree-Search DFS vs. Graph-Search DFS

Tree-Search DFS

- We covered Tree-Search DFS in this course
- Can **be used** for iterative deepening dfs
- Checks new states against those on the **path** from the **root** to the **current node** to avoid infinite loops in finite state space.
- Does not avoid the proliferation of redundant paths.
- Both are non-optimal
- The space complexity can be $O(bm)$

Graph-Search DFS

- You probably covered Graph-Search in COMP2521
- Can **not** be used for iterative deepening dfs!
- Keeps a record of every visited, expanded node
- Both are non-optimal
- The space complexity is as bad as BFS: $O(b^m)$



UNSW
SYDNEY

Depth-Limited Search



UNSW
SYDNEY

Depth-Limited Search

- A depth-first search with a predetermined depth limit ℓ .
- Nodes at depth ℓ , are treated as if they have no successors



UNSW
SYDNEY

Depth-Limited Search

```
function DEPTH-LIMITED-SEARCH (problem, limit) returns a solution, or failure / cutoff
    return Recursive-DLS(Make-Node(problem.INITIAL-STATE), problem, limit)
function RECURSIVE-DLS (node, problem, limit) returns a solution, or failure / cutoff
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    else if limit = 0 return cutoff
    else
        cutoff_occurred?  $\leftarrow$  false
        for each action in problem.ACTIONS(node.STATE) do
            child  $\leftarrow$  CHILD-NODE(problem, node, action)
            result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)
            if result = cutoff then cutoff_occurred?  $\leftarrow$  true
            else if result  $\neq$  failure then return result
        if cutoff_occurred? Then return cutoff else return failure
```

- The algorithm above does not check the current node against those on the path from the root to the current node.

Depth-Limited Search

- Will it cause any issue if we do not check for a loop in depth limited search?



UNSW
SYDNEY

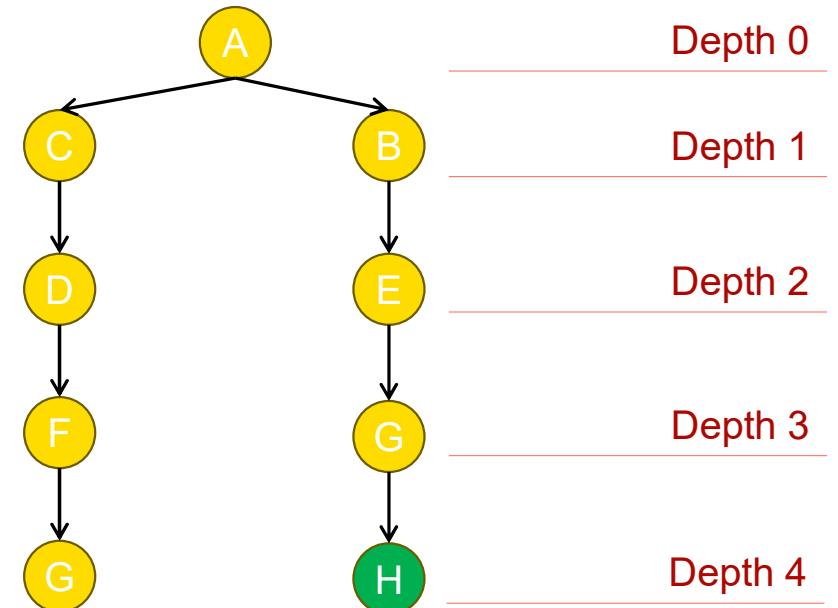
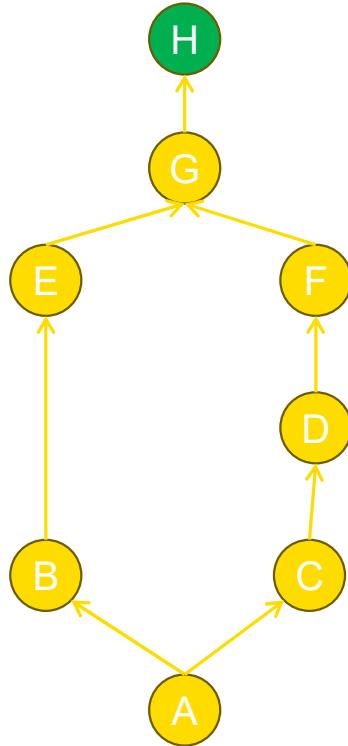
Depth-Limited Search

- Depth limited search can be implemented as a simple recursive algorithm.
- Notice that depth-limited search can be terminate with two kinds of failure: the standard failure value indicates no solution; the cutoff value indicates no solution *within* the depth limit
- The algorithm on previous slide does not check the current node against those on the path from the root to the current node.



UNSW
SYDNEY

How Depth Limited DFS finds a path to goal with depth limit of 4



**Any Question on Comp9414
Armin's Lecture**



<https://forms.office.com/r/bBb6bt2sv7>



UNSW
SYDNEY

Iterative Deepening Depth-First Search (IDDFS)



UNSW
SYDNEY

Iterative Deepening Search

Depth-bounded search: hard to decide on a depth bound

Iterative deepening: Try all possible depth bounds in turn

Combines benefits of depth-first and breadth-first search



Iterative Deepening Search

Tries to combine the benefits of depth-first (low memory) and breadth-first (optimal and complete)

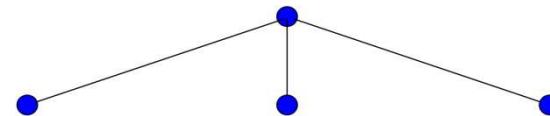
Does a series of depth-limited depth-first searches to depth 1, 2, 3, etc.

Early states will be expanded multiple times, but that might not matter too much because most of the nodes are near the leaves.



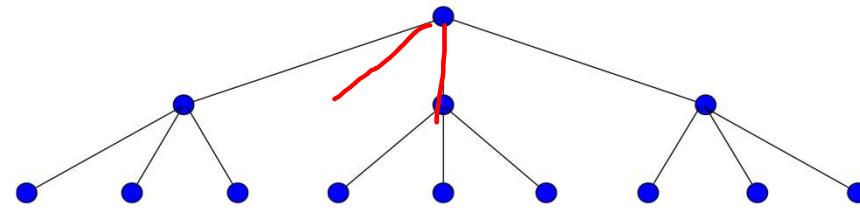
UNSW
SYDNEY

Iterative Deepening Search



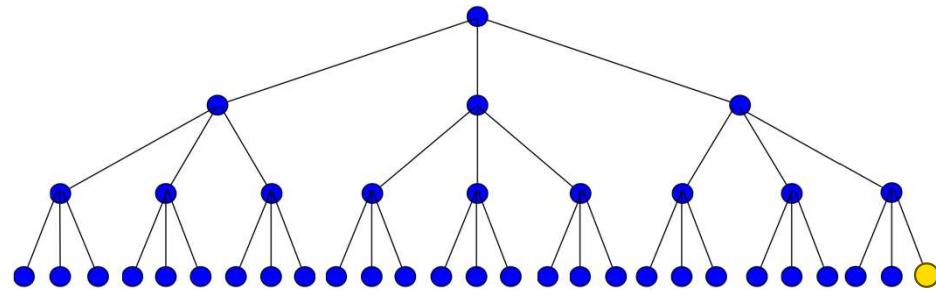
UNSW
SYDNEY

Iterative Deepening Search



UNSW
SYDNEY

Iterative Deepening Search



UNSW
SYDNEY

Properties of Iterative Deepening Search

Complete? Yes.

Time: $O(b^d)$

Space? $O(bd)$

Optimal? Yes, if step costs are identical.

In general, iterative deepening is the preferred search strategy for a large search space where depth of solution is not known



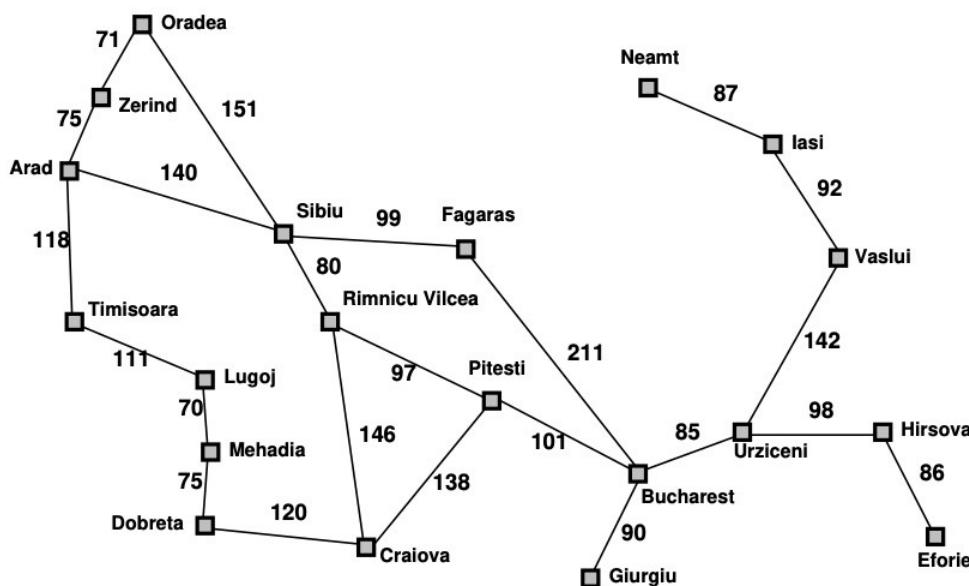
UNSW
SYDNEY

Exercise

This exercise uses the route-finding example with the Romanian map from Russell & Norvig (Artificial Intelligence: A Modern Approach) as an example.

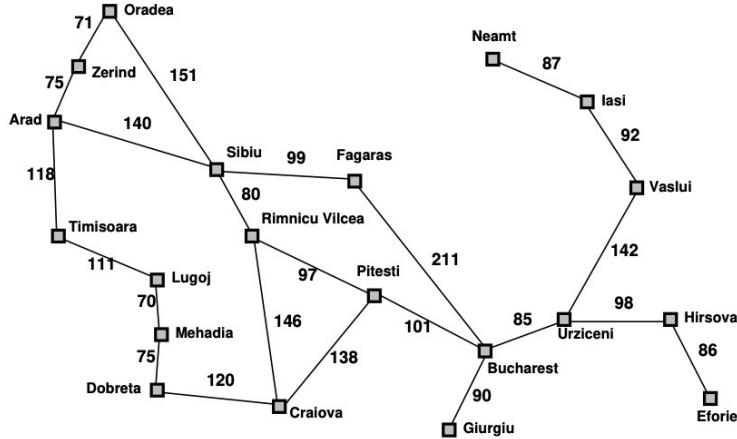
For the route from Arad to Bucharest, what order are nodes in the state space expanded for each of the following algorithms when searching for the shortest path between Arad and Bucharest? Where there is a choice of nodes take the first one by alphabetical ordering. For depth-first search use cycle checking along a path to avoid repeated states that may lead to infinite branches. Stop the search when the goal state is expanded.

* Iterative Deepening
Depth First Search



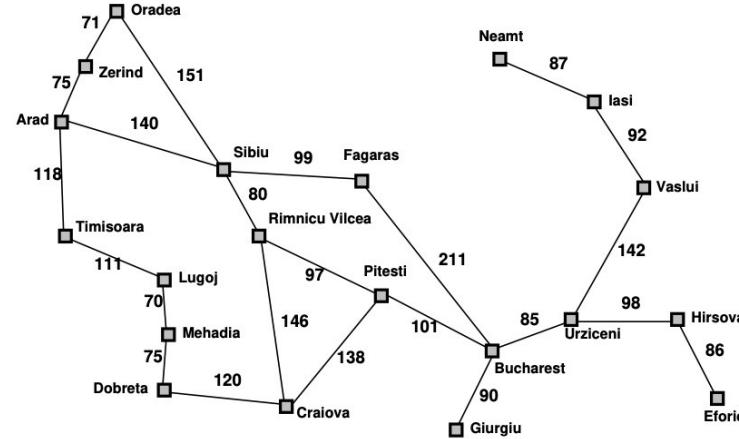
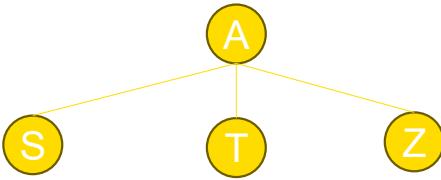
Exercise IDDFS d=1

A

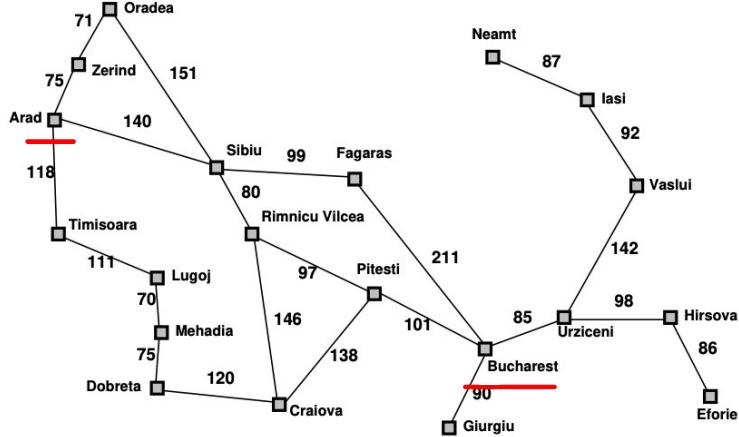
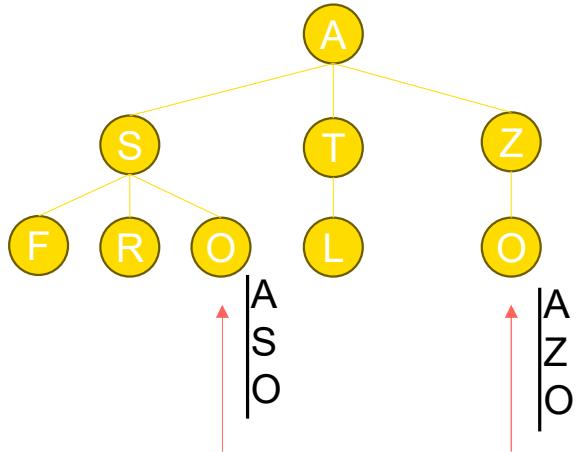


Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

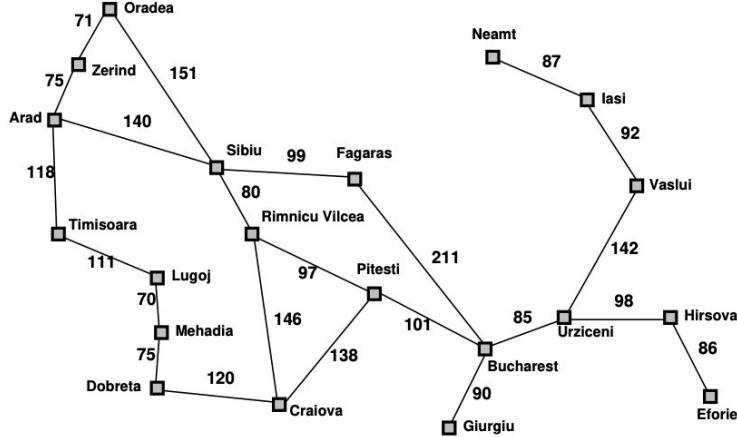
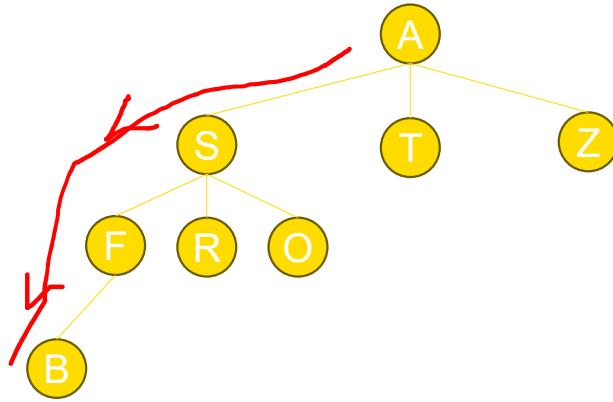
Exercise IDDFS d=2



Exercise IDDFS d=3



Exercise IDDFS d=4



Did you understand IDDFS and DFS?

Did you understand IDDFS and
DFS



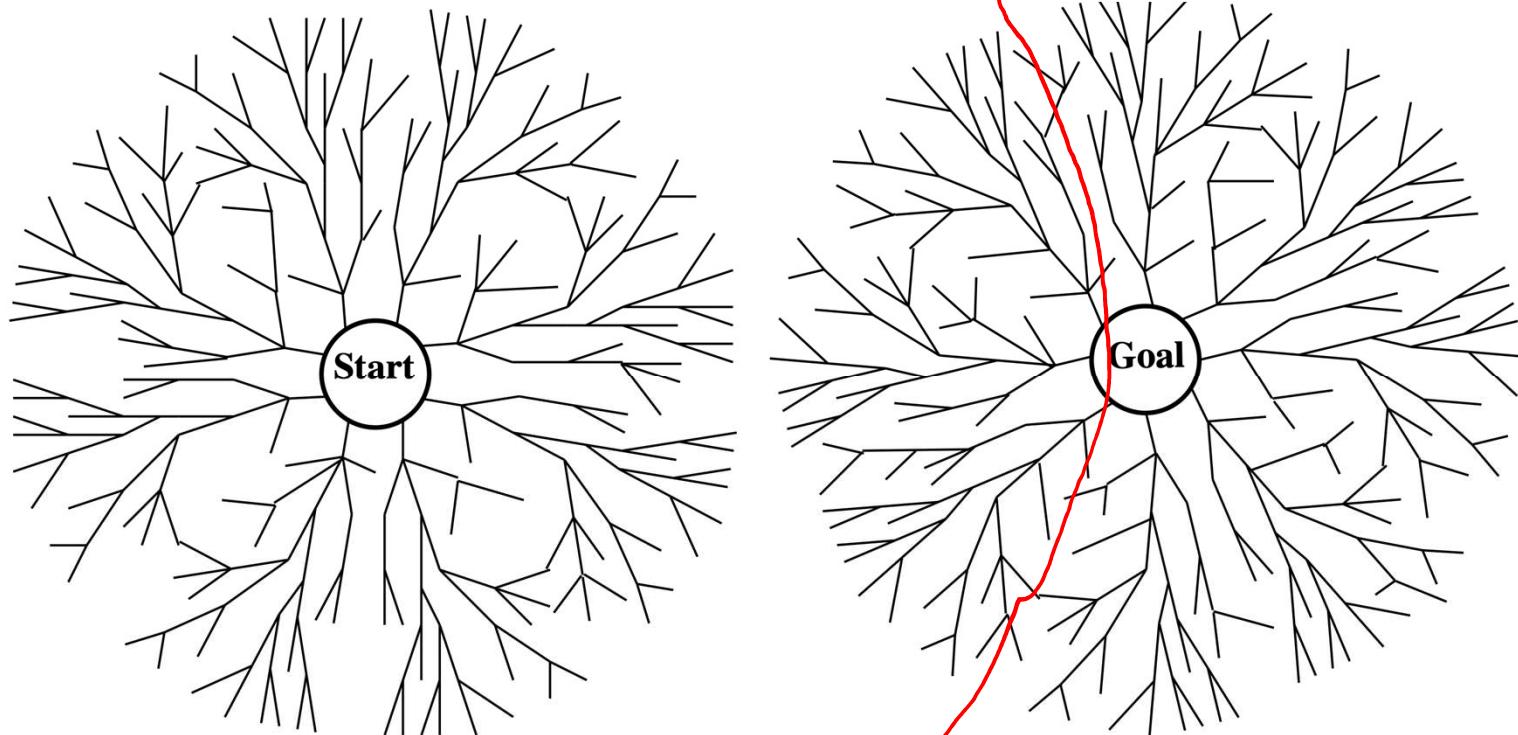
UNSW
SYDNEY

Bidirectional Search



UNSW
SYDNEY

Bidirectional Search



UNSW
SYDNEY

Bidirectional Search

- Search both forward from the initial state and backward from the goal
 - stop when the two searches meet in the middle.
- Need efficient way to check if a new node appears in the other half of the search.
 - Complexity analysis assumes this can be done in constant time, using a hash table.
- Assume branching factor = b in both directions and that there is a solution at depth = d :
 - Then bidirectional search finds a solution in $O(2b^{d/2}) = O(b^{d/2})$ time steps.



UNSW
SYDNEY

Bidirectional Search Analysis

If solution exists at depth d then bidirectional search requires time

$$O(2b^{\frac{d}{2}}) = O(b^{\frac{d}{2}})$$

(assuming constant time checking of intersection)

To check for intersection must have all states from one of the searches in memory, therefore space complexity is $O(b^{\frac{d}{2}})$



UNSW
SYDNEY

Uniform Cost Search (UCS)



UNSW
SYDNEY

Uniform-Cost Search

Lowest-cost-first Search

Sometimes transitions have a cost

Cost of a path is the sum of the costs of its arcs:

$$cost(\langle n_0, \dots, n_k \rangle) = \sum_{i=1}^k cost(\langle n_{i-1}, n_i \rangle)$$

An optimal solution has minimum cost

Delivery robot example:

- cost of arc may be resources (e.g., time, energy) required to execute action represented by the arc
- aim is to reach goal using least resources



UNSW
SYDNEY

Uniform-Cost Search

The simplest search method that is guaranteed to find a minimum cost path is **lowest-cost-first** search or **uniform-cost search**

- similar to breadth-first search, but instead of expanding path with least number of arcs, select path with lowest cost
- implemented by treating the frontier as a priority queue ordered by the [cost function](#)

$$cost(\langle n_0, \dots, n_k \rangle) = \sum_{i=1}^k cost(\langle n_{i-1}, n_i \rangle)$$



Uniform-Cost Search for Delivery Robot

Edges are labelled with cost

- e.g. distance to travel

Sort queue by increasing cost of path to the node

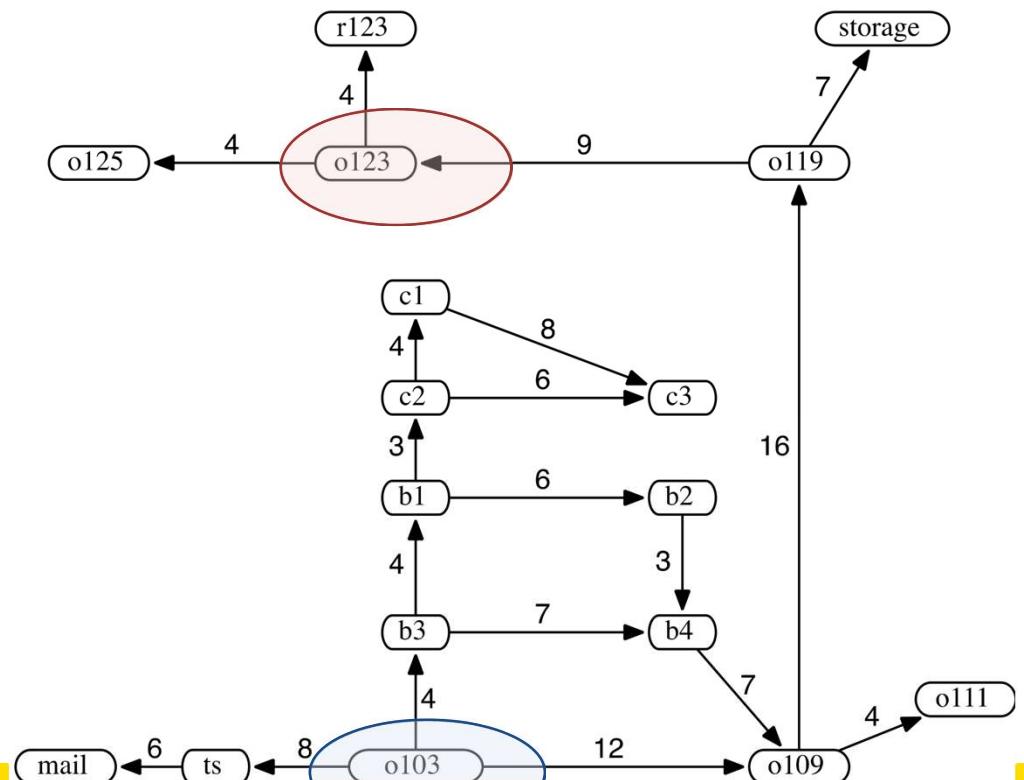
: $[o103_0]$

$[b3_4, ts_8, o109_{12}]$

$[b1_8, ts_8, b4_{11}, o109_{12}]$

$[ts_8, c2_{11}, b4_{11}, o109_{12}, b2_{14}]$

$[c2_{11}, b4_{11}, o109_{12}, mail_{14}, b2_{14}]$



Uniform-Cost Search

- Expand root first, then expand least-cost unexpanded node
- Implementation with priority queue
 - insert nodes in order of increasing path cost - lowest path cost is $\underline{g(n)}$.
- Reduces to breadth-first search when all actions have same cost
- Finds the cheapest goal provided path cost is monotonically increasing along each path (i.e. no negative-cost steps)

Properties of Uniform-Cost Search

Complete? Yes, if b is finite and if transition $\text{cost} \geq \epsilon$ with $\epsilon > 0$

Time? Worst case, $O(b^{\lceil C^*/\epsilon \rceil})$ where C^* = cost of the optimal solution
every transition costs at least ϵ
 \therefore cost per step is $\frac{C^*}{\epsilon}$

Space? $O(b^{\lceil C^*/\epsilon \rceil})$, $b^{\lceil C^*/\epsilon \rceil} = b^d$ if all step costs are equal

Optimal? Yes – nodes expanded in increasing order of lower path cost, $g(n)$



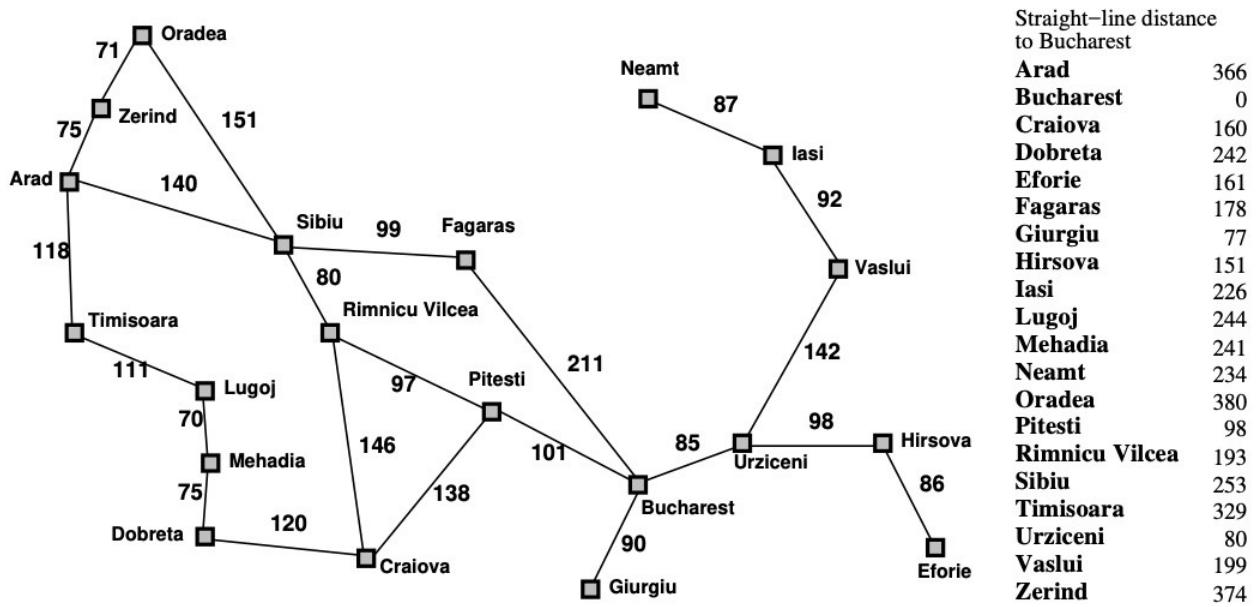
UNSW
SYDNEY

Exercise

This exercise uses the route-finding example with the Romanian map from Russell & Norvig (Artificial Intelligence: A Modern Approach) as an example.

For the route from Arad to Bucharest, what order are nodes in the state space expanded for each of the following algorithms when searching for the shortest path between Arad and Bucharest? Where there is a choice of nodes take the first one by alphabetical ordering. For uniform0cost search include a check that nodes with the same state as previously expanded nodes are not added to the frontier. Stop the search when the goal state is expanded.

* Uniform Cost Search



Complexity Results for Uninformed Search

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Time	$O(b^d)$	$O(b^{\lceil C^*/\varepsilon \rceil})$	$O(b^m)$	$O(b^k)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{\lceil C^*/\varepsilon \rceil})$	$O(bm)$	$O(bk)$	$O(bd)$
Complete?	Yes ¹	Yes ²	No	No ⁴	Yes ¹
Optimal ?	Yes ³	Yes	No	No	Yes ³

b = branching factor, d = depth of the shallowest solution,

m = maximum depth of the search tree, k = depth limit.

1 = complete if b is finite.

2 = complete if b is finite and step costs $\geq \varepsilon$ with $\varepsilon > 0$.

3 = optimal if actions all have the same cost.

4 = incomplete if the goal is not within the depth bound

Any Question on Comp9414
Armin's Lecture



<https://forms.office.com/r/bBb6bt2sv7>



UNSW
SYDNEY



UNSW
SYDNEY

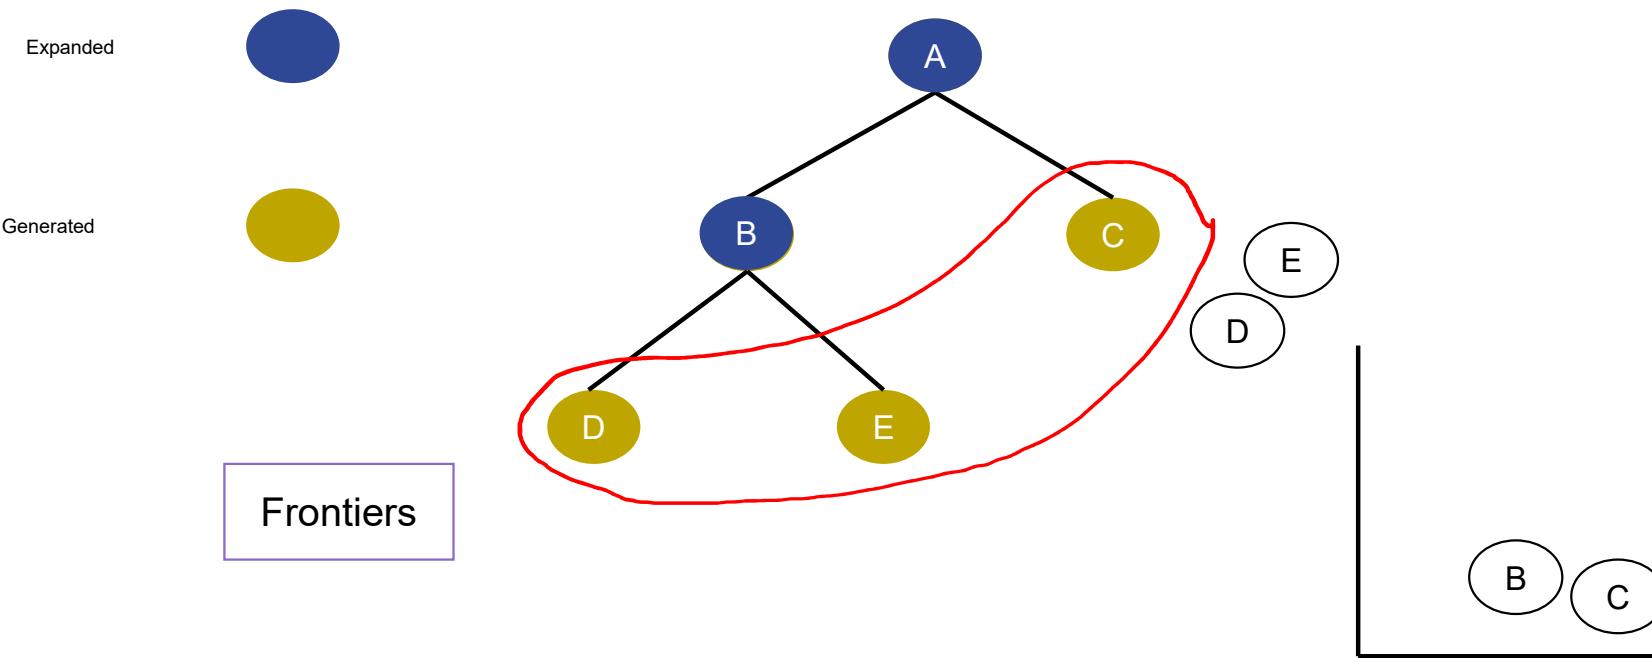
Solving Problems by Searching

Informed Search

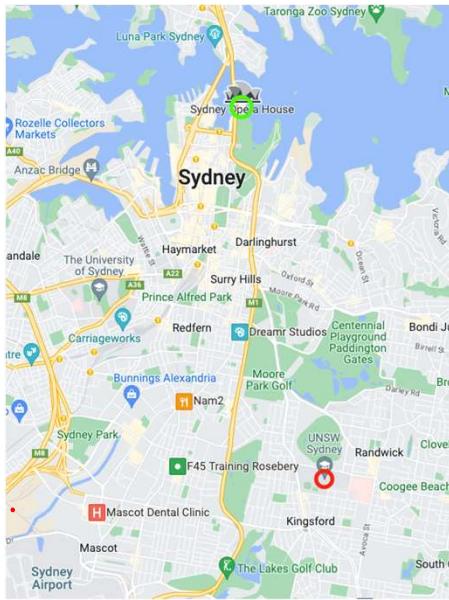


UNSW
SYDNEY

Search strategies



Focus the search on the path that is most likely to lead to the goal state



UNSW
SYDNEY

Overview Informed Search

Uninformed Search	Background & Terminology
	Breath First Search (BFS)
	Depth First Search (DFS)
	Uniform Cost Search (UCS)
	Depth Limited Search
	Iterative Deepening Search (IDDFS)
	Bidirectional Search
Informed Searches	Heuristic
	Greedy Best-First Search
	A* Search



Search Strategies

General Search algorithm:

- add initial state to the list
- repeat:
 - take node from the list
 - test if it is a goal state; if so, terminate
 - “expand” it, i.e. generate successor nodes and add them to the list

Search strategies are distinguished by the order in which new nodes are added to the queue of nodes awaiting expansion.



UNSW
SYDNEY

Uniformed vs Informed Search

- **Uninformed** - keeps searching until it stumbles on goal
 - No domain knowledge
- **Informed** - searches in direction of best guess to goal
 - Uses domain knowledge



UNSW
SYDNEY

Informed (Heuristic) Search

- Informed search strategy
 - use problem-specific knowledge more efficiently than uninformed search
- Uninformed search algorithms have no information about problem other than its definition.
 - some can solve any solvable problem, none of them can do it efficiently
- Informed search algorithms can do well given guidance on where to look for solutions.
- Implemented using a **priority queue** to store frontier nodes

Search Strategies

BFS and DFS treat all new nodes the same way:

- BFS add all new nodes to the back of the list
- DFS add all new nodes to the front of the list

Best First Search uses an evaluation function $f()$ to order the nodes in the list

- Similar to uniform cost search

Informed or Heuristic:

- Greedy Search $f(n) = h(n)$ (estimates cost from node n to goal)
- A* Search $f(n) = g(n) + h(n)$ (cost from start to n plus estimated cost to goal)



UNSW
SYDNEY

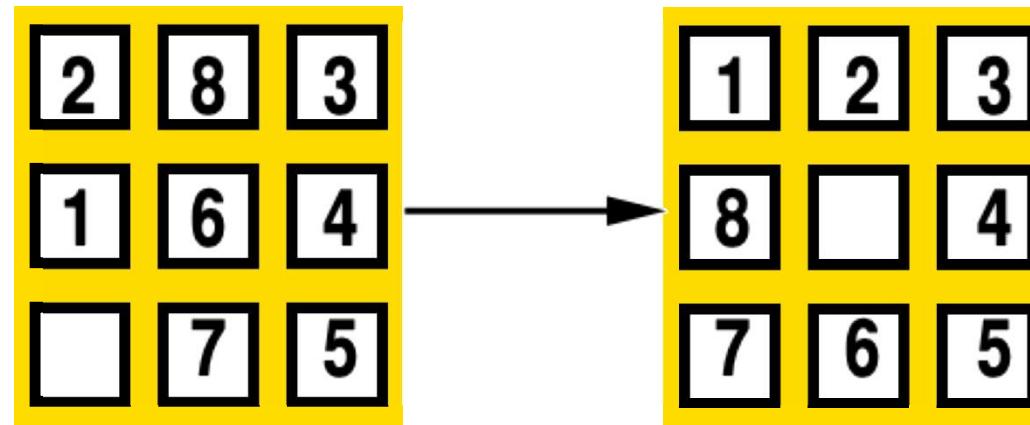
Heuristics

- Heuristics are “rules of thumb” for deciding which alternative is best
- Heuristic must ***underestimate*** actual cost to get from current node to goal
 - Called an **admissible heuristic**
- Denoted ***h(n)***
 - $h(n) = 0$ whenever n is a **goal** node



Heuristics — Example

8-Puzzle — number of tiles out of place



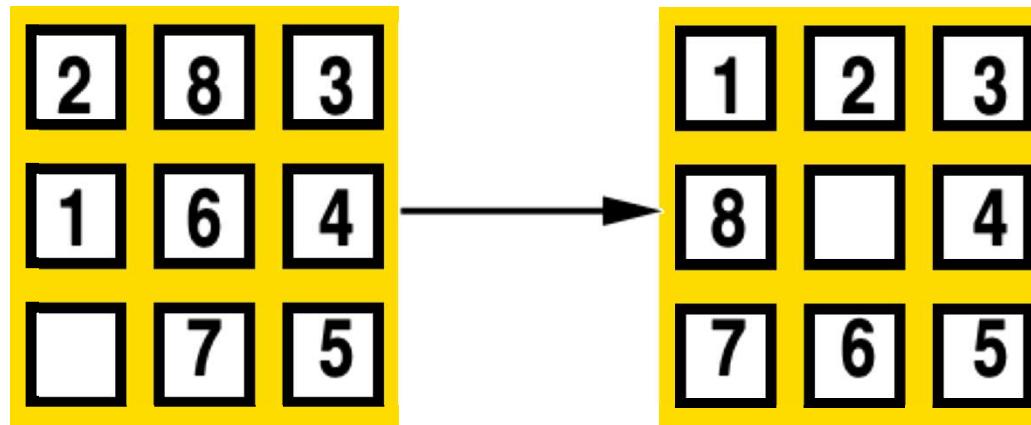
$$h(n) = 5$$



UNSW
SYDNEY

Heuristics — Example

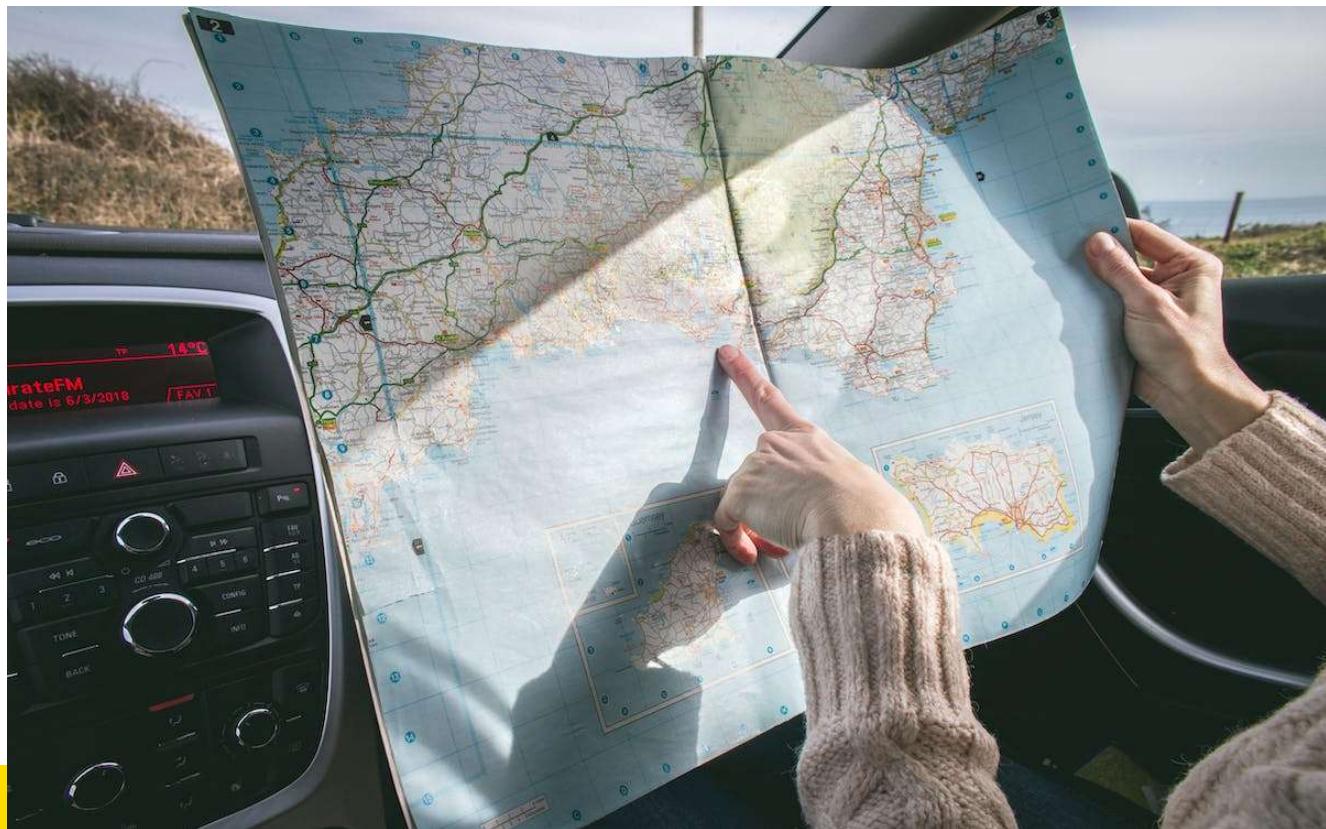
8-Puzzle — Manhattan distance (distance tile is out of place)



$$h(n) = 1 + 1 + 0 + 0 + 0 + 1 + 1 + 2 = 6$$



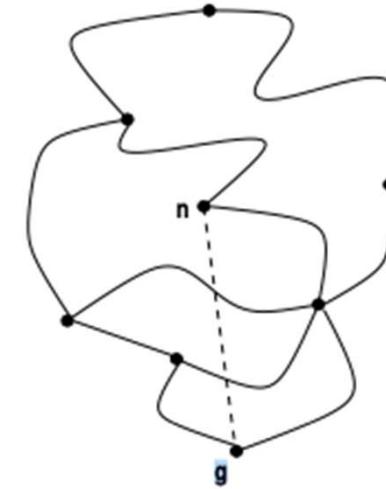
Heuristics — Example



UNSW
SYDNEY

Heuristics — Example

Another common heuristic is the straight-line distance (“as the crow flies”) from node to goal



Therefore $h(n) = \text{distance from } n \text{ to } g$

Example Heuristic Functions

If nodes are points on a Euclidean plane and cost is distance, $h(n)$ can be straight-line distance from n to closest goal.

If nodes are locations and cost is time,
can use distance to goal divided by maximum speed.

- If goal is to collect a bunch of coins and not run out of fuel, cost is an estimate of how many steps to collect rest of the coins, refuel when necessary, and return to goal.

Heuristic function can be found by simplifying calculation of true cost

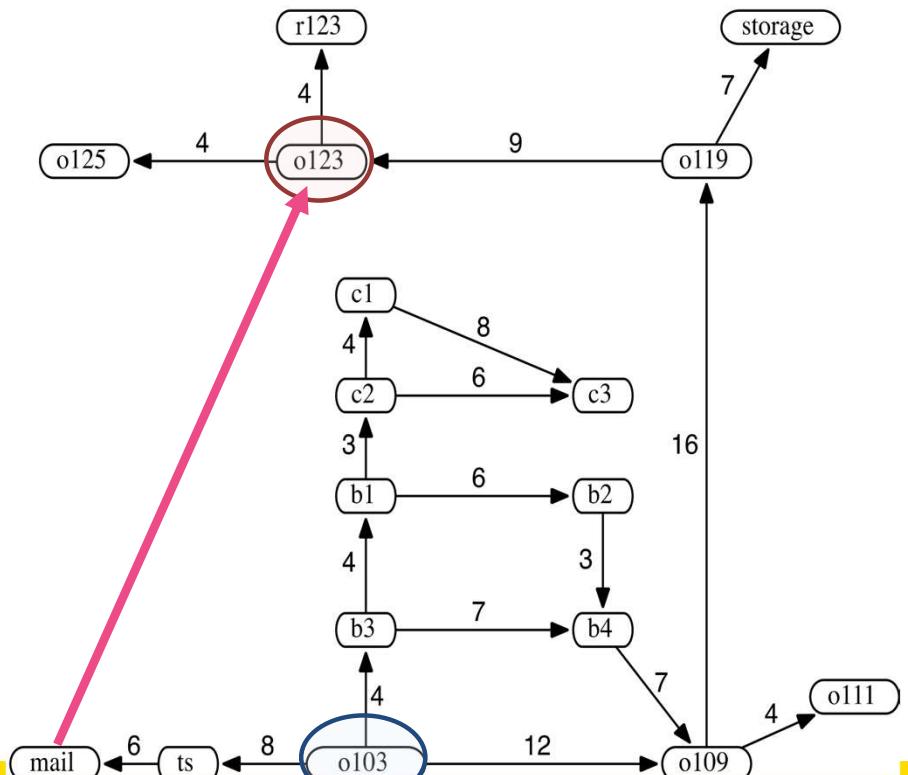
Delivery Robot Heuristic Function

Use straight-line distance as heuristic, and assume these values:

$h(\text{mail}) = 26$	$h(ts) = 23$	$h(o103) = 21$
$h(o109) = 24$	$h(o111) = 27$	$h(o119) = 11$
$h(o123) = 4$	$h(o125) = 6$	$h(r123) = 0$
$h(b1) = 13$	$h(b2) = 15$	$h(b3) = 17$
$h(b4) = 18$	$h(c1) = 6$	$h(c2) = 10$
$h(c3) = 12$	$h(\text{storage}) = 12$	

$h(\text{loc})$ = distance from loc to goal

Heuristic function can be extended to paths by making heuristic value of path equal to heuristic value of node at the end of the path: $h(\langle n_0, \dots, n_k \rangle) = h(n_k)$.



Exercise

What are some other heuristics that you can think of?



UNSW
SYDNEY

Greedy Best-First Search



UNSW
SYDNEY

Greedy Best-First Search

Always select node closest to goal according to heuristic function

$h(n)$ is estimated cost to goal

- $h(n) = 0$ if n is a goal state

Frontier is a **priority queue** ordered by h .

“Greedy” algorithm takes “best” node first.

$$f(n) = h(n)$$

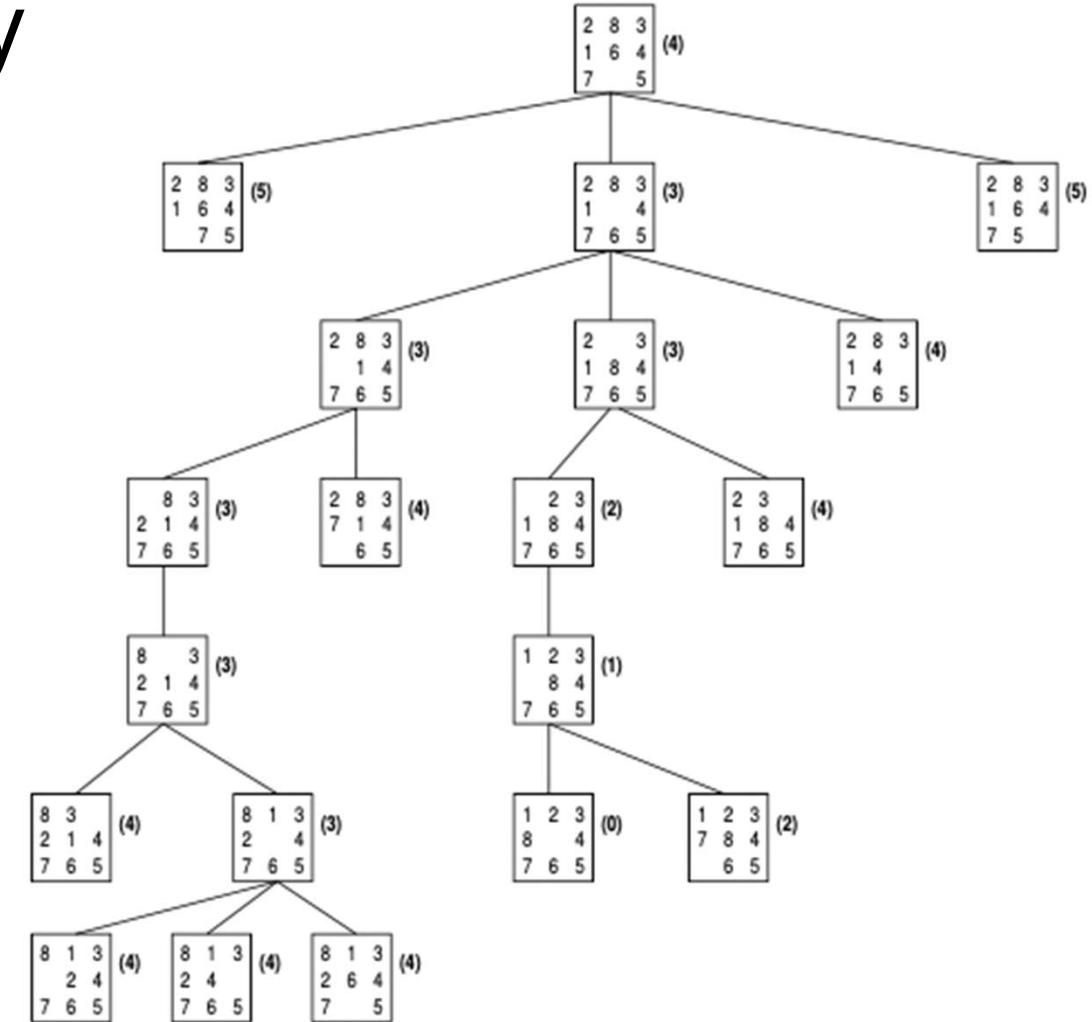
Like depth-first search, except pick next node by $h(n)$



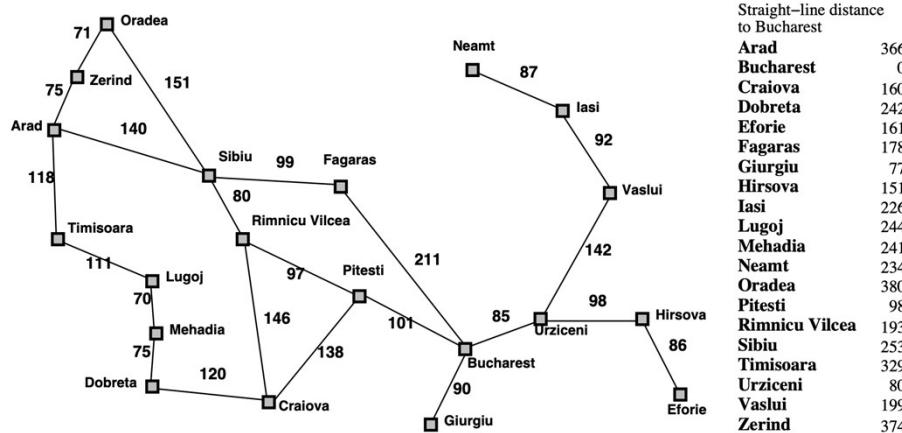
UNSW
SYDNEY

Examples of Greedy Best-First Search

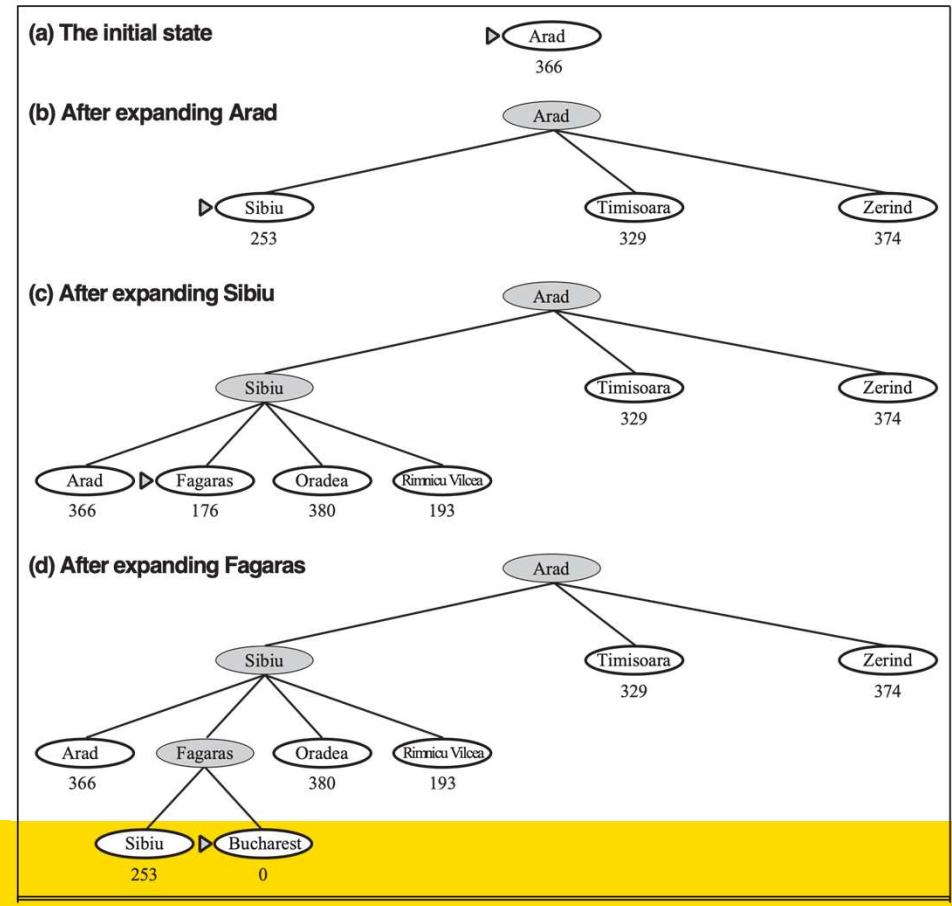
h is number of tiles out of place



Examples of Greedy Best-First Search



- Stages in a greedy best-first tree search for route from Arad to Bucharest with the straight-line distance heuristic.
- Note that **straight-line distances are less than actual distances** in map.

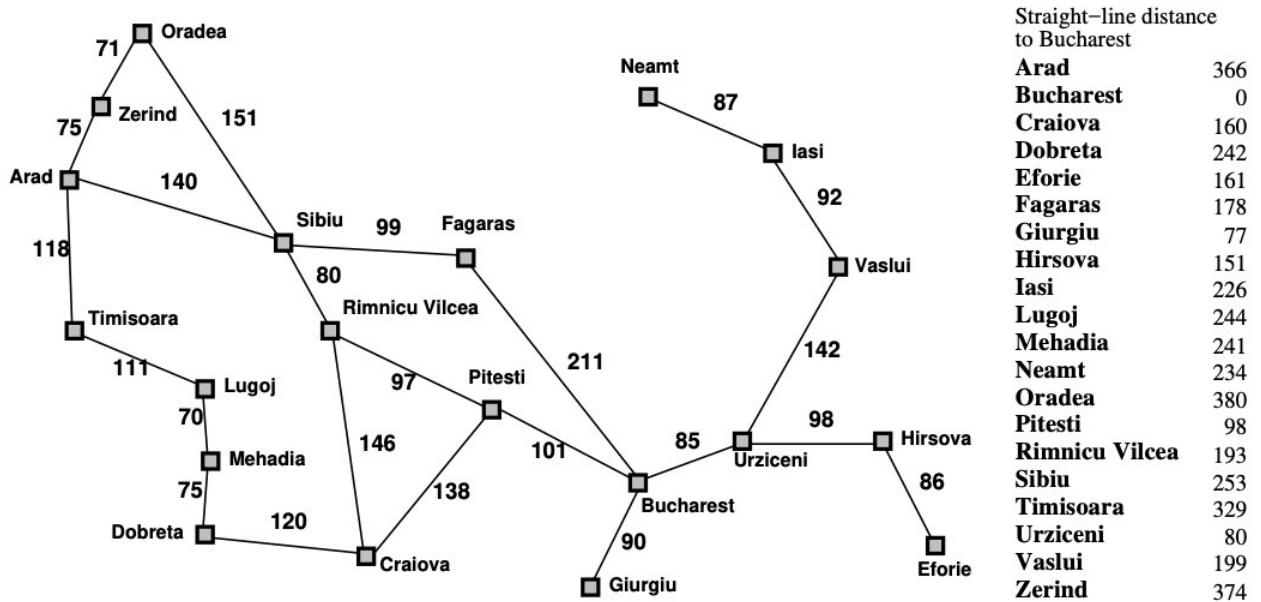


Exercise

This exercise uses the route-finding example with the Romanian map from Russell & Norvig (Artificial Intelligence: A Modern Approach) as an example.

For the route from Arad to Bucharest, what order are nodes in the state space expanded for each of the following algorithms when searching for the shortest path between Arad and Bucharest? Where there is a choice of nodes take the first one by alphabetical ordering. Stop the search when the goal state is expanded.

* Greedy best-first search



Properties of Greedy Best-First Search

Complete: No. Can get stuck in loops.

(Complete in finite space with repeated-state checking)

Time: $O(b^m)$, where m is the maximum depth in search space.

Space: $O(b^m)$ (retains all nodes in memory)

Optimal: No.

Greedy Search is not efficient

However, a good heuristic can reduce time and memory costs substantially.



UNSW
SYDNEY

A* Search



UNSW
SYDNEY

Uniform-Cost Search



Greedy Search



A* Search

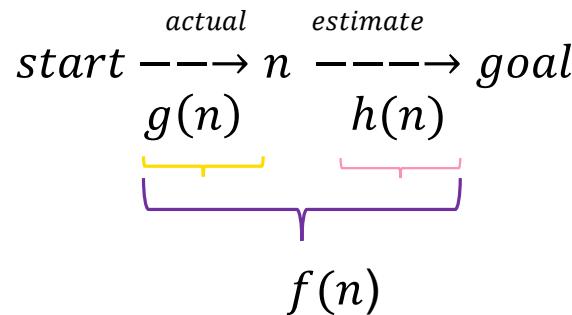


**UNSW
SYDNEY**

A* Search

Use both cost of path generated and estimate to goal to order nodes on the frontier

- $g(n)$ = cost of path from start to n
- $h(n)$ = estimate from n to goal



- Order priority queue using function $f(n) = g(n) + h(n)$
- $f(n)$ is the estimated cost of the cheapest solution extending this path



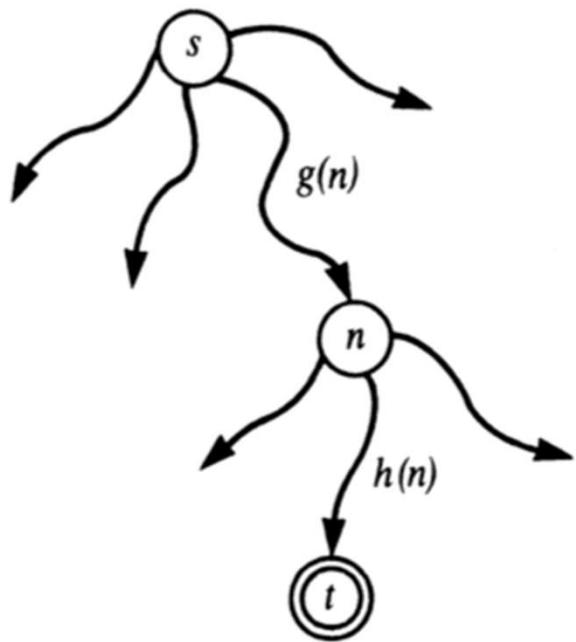
A* Search

- Combines uniform-cost search and greedy search
- Greedy Search minimises $h(n)$
 - efficient but not optimal or complete
- Uniform Cost Search minimises $g(n)$
 - optimal and complete but not efficient



UNSW
SYDNEY

Heuristic Function



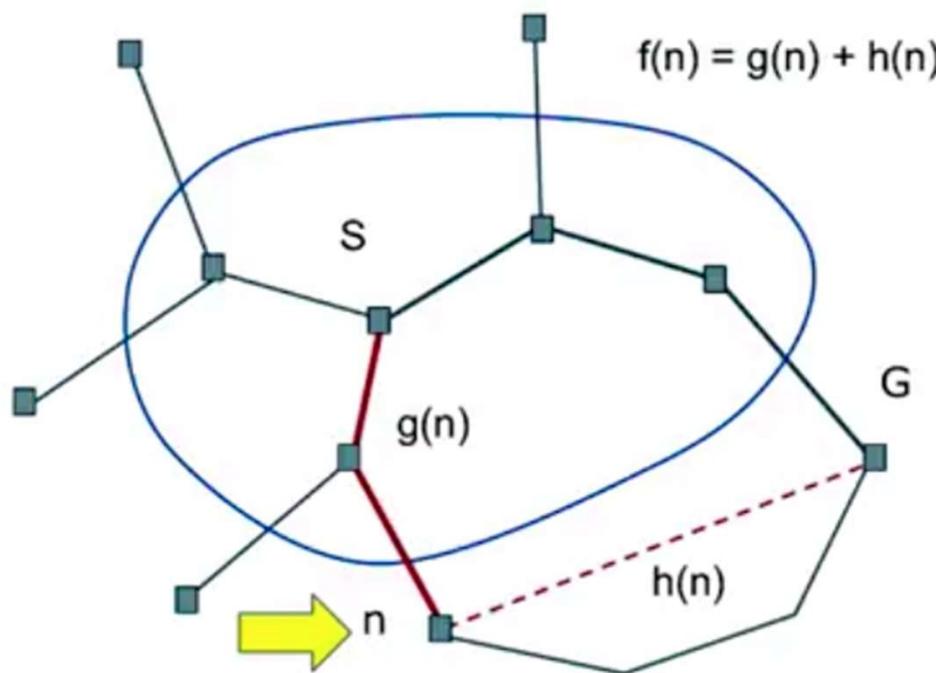
Heuristic estimate $f(n) = \text{cost of the cheapest path from } s \text{ to } t \text{ via } n: f(n) = g(n) + h(n)$

g(n) is the cost the path from s to n

h(n) is an estimate of the cost of an optimal path from n to t.



A* Search



S = start

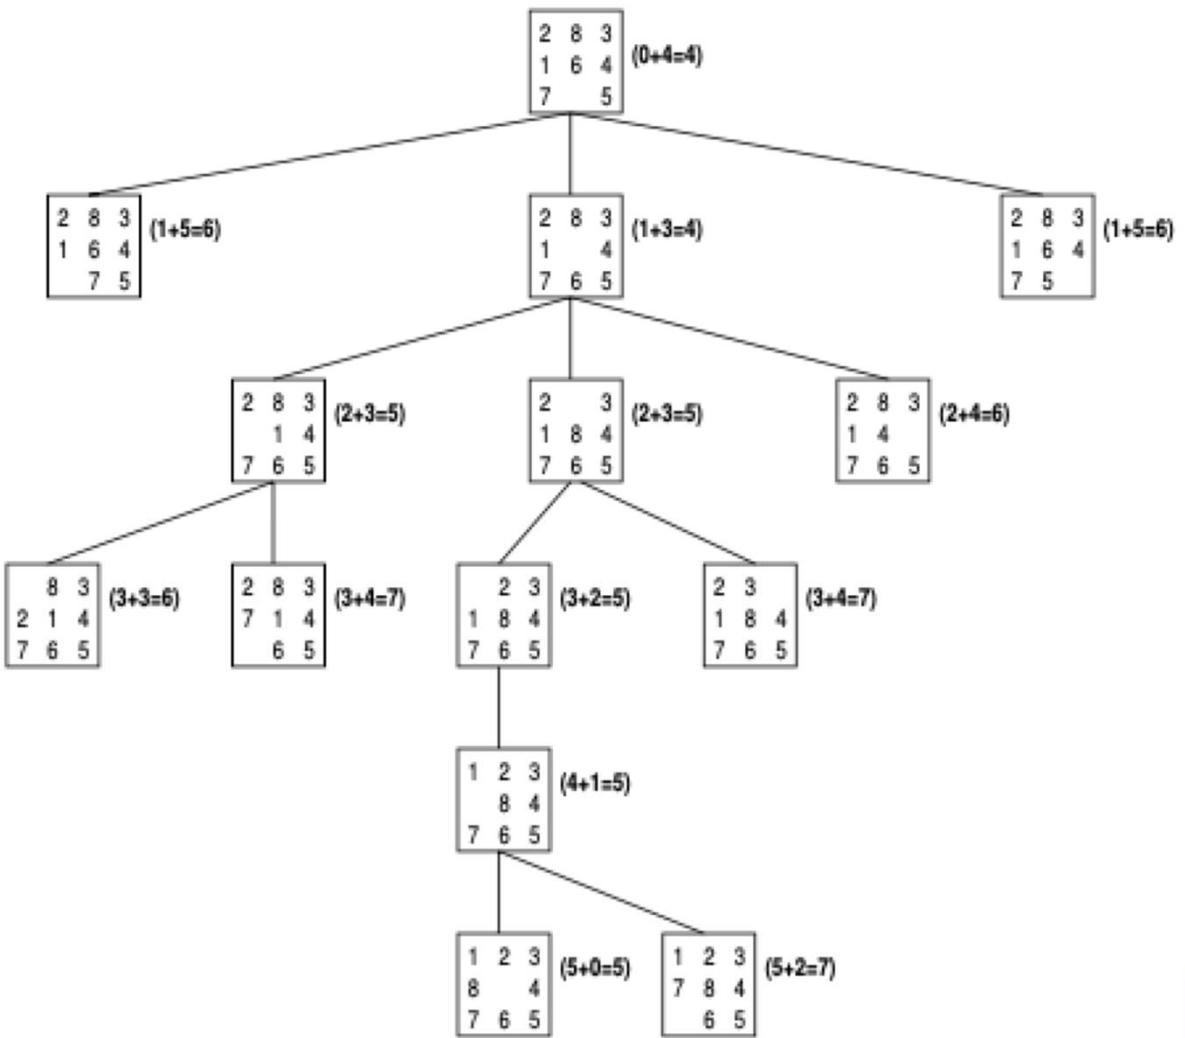
G = Goal

n = current node

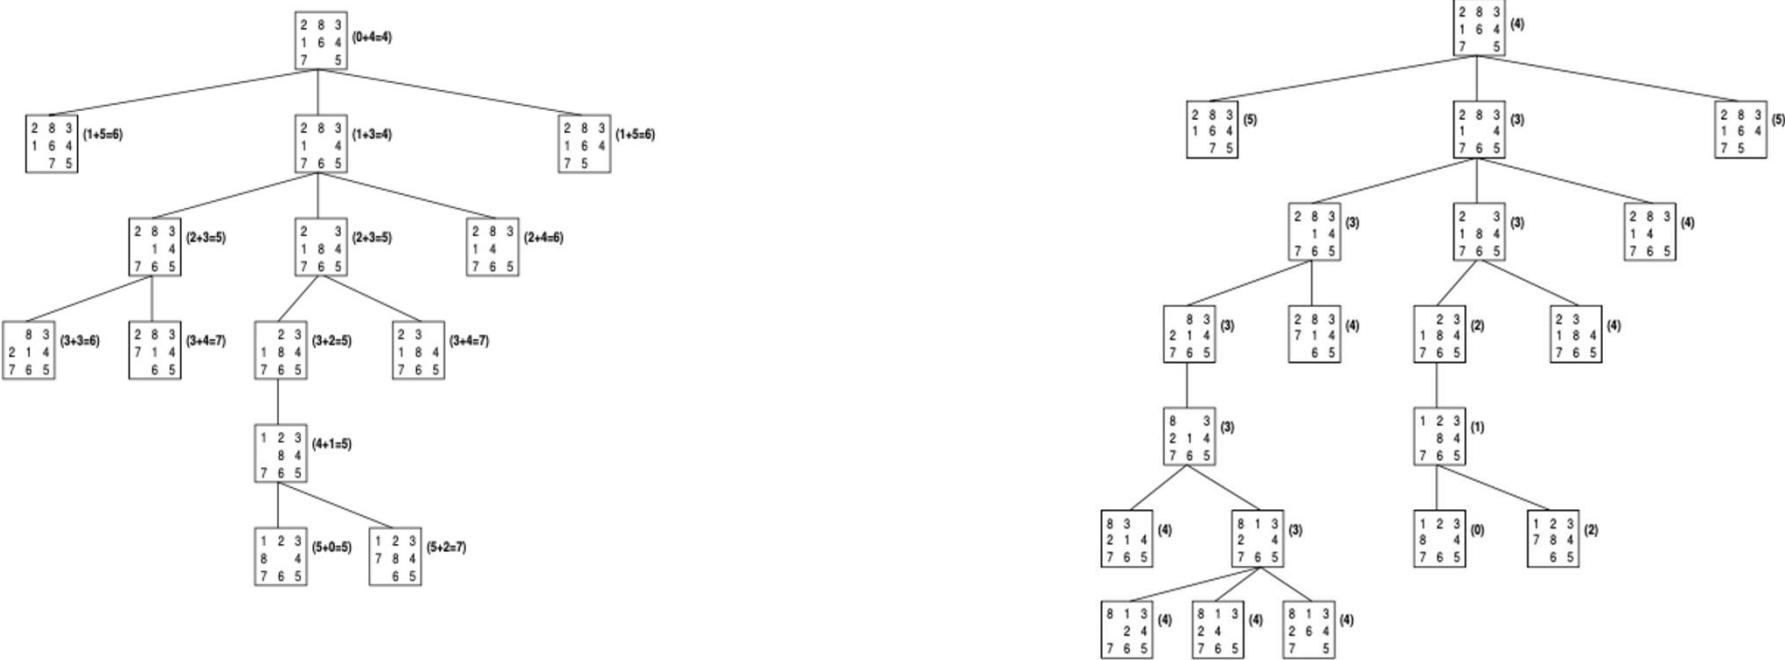
$g(n)$ = actual cost from **S** to **n**

$h(n)$ = estimated distance from **n** to **G**

A* Search



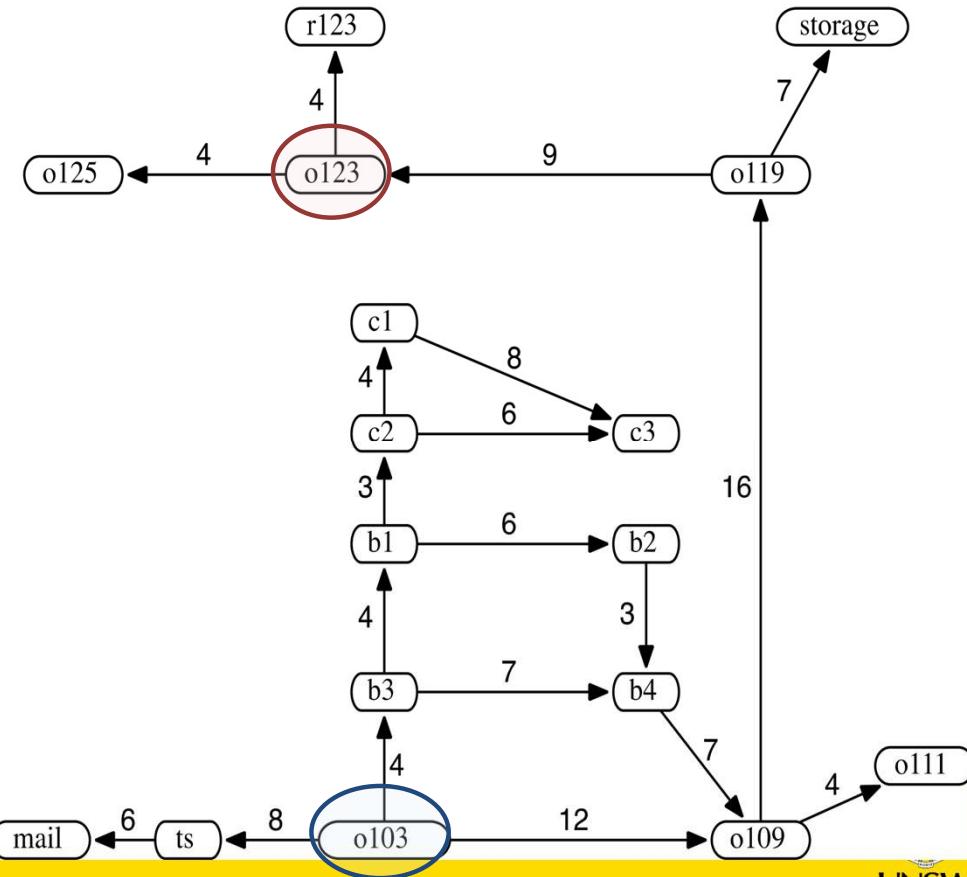
A* vs Greedy



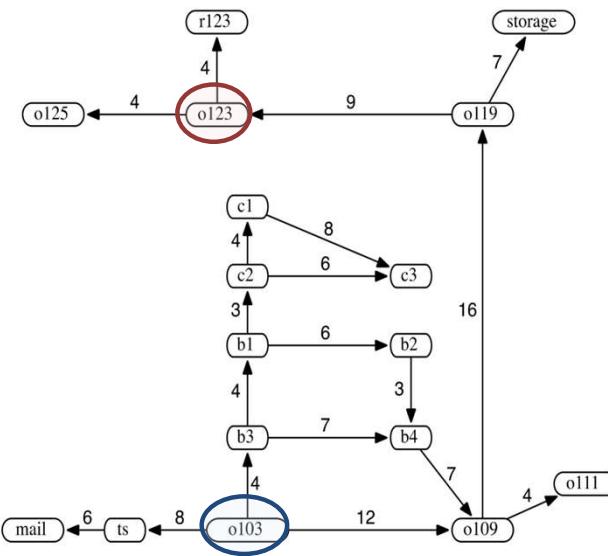
Delivery Robot Heuristic Function

Use straight-line distance as heuristic, and assume these values:

$$\begin{array}{lll}
 h(\text{mail}) = 26 & h(ts) = 23 & h(o103) = 21 \\
 h(o109) = 24 & h(o111) = 27 & h(o119) = 11 \\
 h(o123) = 4 & h(o125) = 6 & h(r123) = 0 \\
 h(b1) = 13 & h(b2) = 15 & h(b3) = 17 \\
 h(b4) = 18 & h(c1) = 6 & h(c2) = 10 \\
 h(c3) = 12 & h(\text{storage}) = 12
 \end{array}$$



A* Search - The Delivery Robot



$$\begin{aligned}
 h(\text{mail}) &= 26 & h(ts) &= 23 & h(o103) &= 21 \\
 h(o109) &= 24 & h(o111) &= 27 & h(o119) &= 11 \\
 h(o123) &= 4 & h(o125) &= 6 & h(r123) &= 0 \\
 h(b1) &= 13 & h(b2) &= 15 & h(b3) &= 17 \\
 h(b4) &= 18 & h(c1) &= 6 & h(c2) &= 10 \\
 h(c3) &= 12 & h(\text{storage}) &= 12
 \end{aligned}$$

1. [o103₂₁] $h(o103) = 21$
 2. [b3₂₁, ts₃₁, o109₃₆] $f((o103, b3)) = g((o103, b3)) + h(b3) = 4 + 17 = 21$
 3. [b1₂₁, b4₂₉, ts₃₁, o109₃₆]
 4. [c2₂₁, b2₂₉, b4₂₉, ts₃₁, o109₃₆]
 5. [c1₂₁, b2₂₉, b4₂₉, c3₂₉, ts₃₁, o109₃₆]
 6. [b2₂₉, b4₂₉, c3₂₉, ts₃₁, c3₃₅, o109₃₆]
 7. [b4₂₉, ts₃₁, c3₃₅, o109₃₆]
 8. [ts₃₁, c3₃₅, b4₃₅, o109₃₆, o109₄₂]
-

$$h(o103) = 21$$

- Lowest-cost path is eventually found.
- Forced to try many different paths, because some temporarily seem to have the lowest cost.
- Still does better than lowest-cost-first search and greedy best-first search.



Optimality of A*

- Heuristic h is said to be **admissible** if
$$\forall n h(n) \leq h^*(n)$$
 where $h^*(n)$ is the true cost from n to goal
- If h is **admissible** then $f(n)$ never overestimates the actual cost of the best solution through n .
$$f(n) = g(n) + h(n),$$
 where $g(n)$ is the actual cost to n and $h(n)$ is an underestimate
- Example: $h = \text{straight line distance}$ is admissible because the shortest path between any two points is a line.
- A^* is optimal if h is **admissible**.
- Admissible heuristics are by nature optimistic because they think the cost of solving the problem is less than it actually is.



UNSW
SYDNEY

Optimality of A* Search

Complete: Yes, unless infinitely many nodes with $f \leq \text{cost of solution}$

Time: Exponential in $\text{relative error in } h \times \text{length of solution}$

Space: Keeps all expanded nodes in memory

Optimal: Yes (assuming h is admissible).

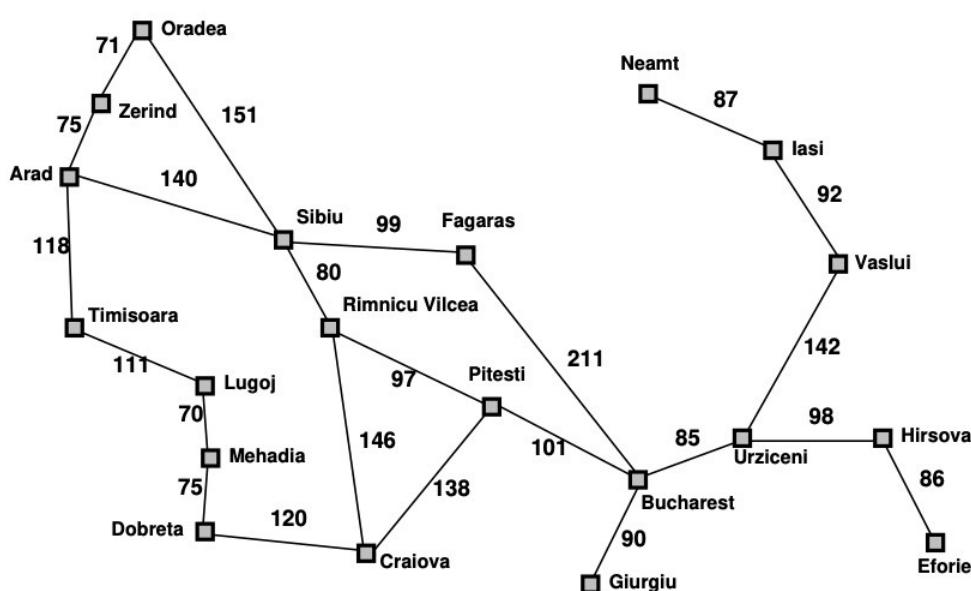


Exercise

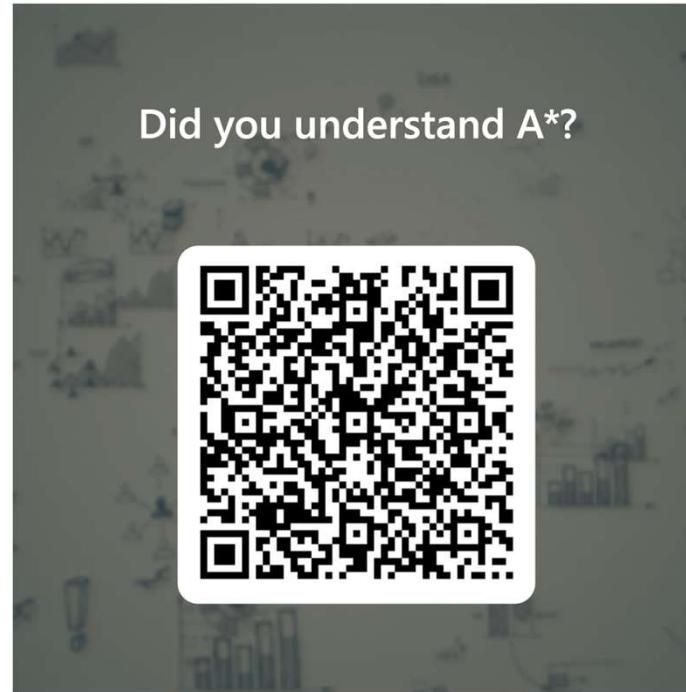
This exercise uses the route-finding example with the Romanian map from Russell & Norvig (Artificial Intelligence: A Modern Approach) as an example.

For the route from Arad to Bucharest, what order are nodes in the state space expanded for each of the following algorithms when searching for the shortest path between Arad and Bucharest? Where there is a choice of nodes take the first one by alphabetical ordering. For uniform0cost search include a check that nodes with the same state as previously expanded nodes are not added to the frontier. Stop the search when the goal state is expanded.

- A*



Did you understand A*?



UNSW
SYDNEY

homework

What sort of search will greedy search emulate if we run it with:

- $h(n) = -g(n)$?
- $h(n) = g(n)$
- $h(n) = \text{number of steps from initial state to node } n$?



UNSW
SYDNEY

Iterative Deepening A* Search



UNSW
SYDNEY

Iterative Deepening A* Search

Iterative Deepening A* is a low-memory variant of A* that performs a series of depth-first searches but cuts off each search when the f exceeds current threshold, initially $f(start)$.

The threshold is increased with each successive search.



Homework

If $h(n)$ is an underestimate of the actual cost, then it is called admissible heuristic. An A* search with an admissible heuristic always finds an optimal path. Can you show why this is always the case?

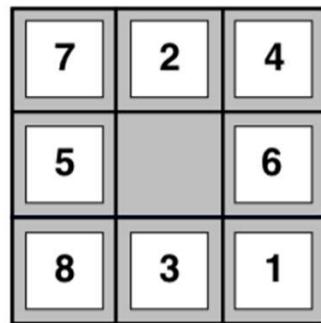


UNSW
SYDNEY

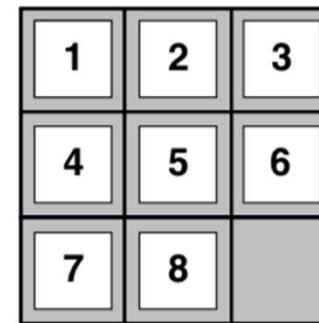
Examples of Admissible Heuristics

$h_1(n)$ = total number of misplaced tiles

$h_2(n)$ = total Manhattan distance = \sum distance from goal position



Start State



Goal State

$$h_1(\text{Start}) = 6$$

$$h_2(\text{Start}) = 4 + 0 + 3 + 3 + 1 + 0 + 2 + 1 = 14$$



UNSW
SYDNEY

More on Heuristic



UNSW
SYDNEY

How to Find Heuristic Functions ?

Admissible heuristics can often be derived from the exact solution cost of a simplified or “relaxed” version of the problem.
(i.e. with some of the constraints weakened or removed)

- If the rules of the 8-puzzle are relaxed so that a tile can move anywhere, then $h_1(n)$ gives the shortest solution.
- If the rules are relaxed so that a tile can move to any adjacent square, then $h_2(n)$ gives the shortest solution.



Dominance

if $h_2(n) \geq h_1(n)$ for all n (both admissible)

- h_2 **dominates** h_1 and is better for search.
- Try make the heuristic h as large as possible, without exceeding h^* .

typical search costs:

14-puzzle	IDS	= 3,473,941 nodes
	$A^*(h_1)$	= 539 nodes
	$A^*(h_2)$	= 113 nodes
24-puzzle	IDS	$\approx 54 \times 10^9$ nodes
	$A^*(h_1)$	= 39,135 nodes
	$A^*(h_2)$	= 1,641 nodes

Summary of Informed Search

Informed search makes use of problem-specific knowledge to guide progress of search

This can lead to a significant improvement in performance

Heuristics can be applied to reduce search cost.

Greedy Search tries to minimise cost from current node n to the goal.



Summary of Informed Search

A* combines the advantages of Uniform-Cost Search and Greedy Search

A* is complete, optimal and optimally efficient among all optimal search algorithms.

Memory usage is still a concern for A*. IDA* is a low-memory variant.

Much research has gone into admissible heuristics

- Even on the automatic generation of admissible heuristics



How to solve any problem using search?



UNSW
SYDNEY

How to solve any problem using search?

1. Turn the states, actions and transition function into a machine-readable format
2. Choose a search strategy according to the problem



UNSW
SYDNEY

Updates

Update	Page	Thanks to
Children of Sibiu are Fagaras, Oradea and Rimincu Vlcea. The way asked in the question was to put them in alphabetical order. It means the order should have been R, O, F . However, on the original slide, they were placed in a wrong order which is not fixed.	Page 27	Amal Ramadan Alshammari



UNSW
SYDNEY

Anonymous feedback for Armin's Lecture

<https://forms.office.com/r/eFhvP4dz0y>



UNSW
SYDNEY